

# **Antialiasing & Compositing**

CS4620 Lecture 25

# Pixel coverage

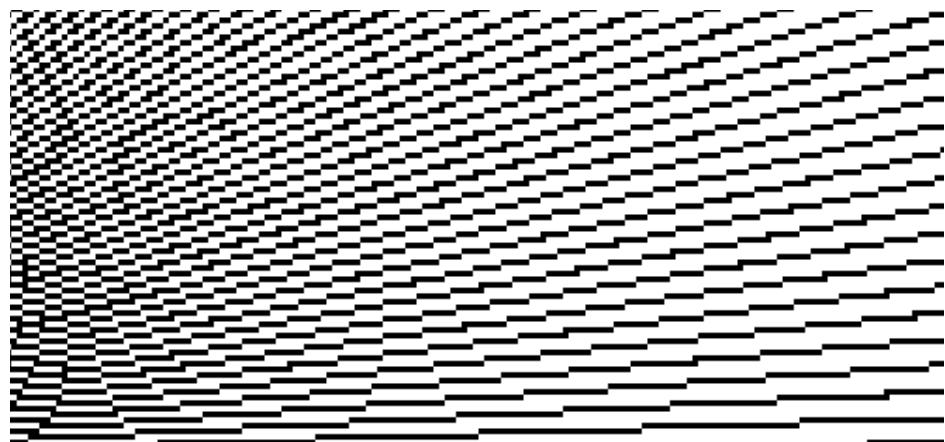
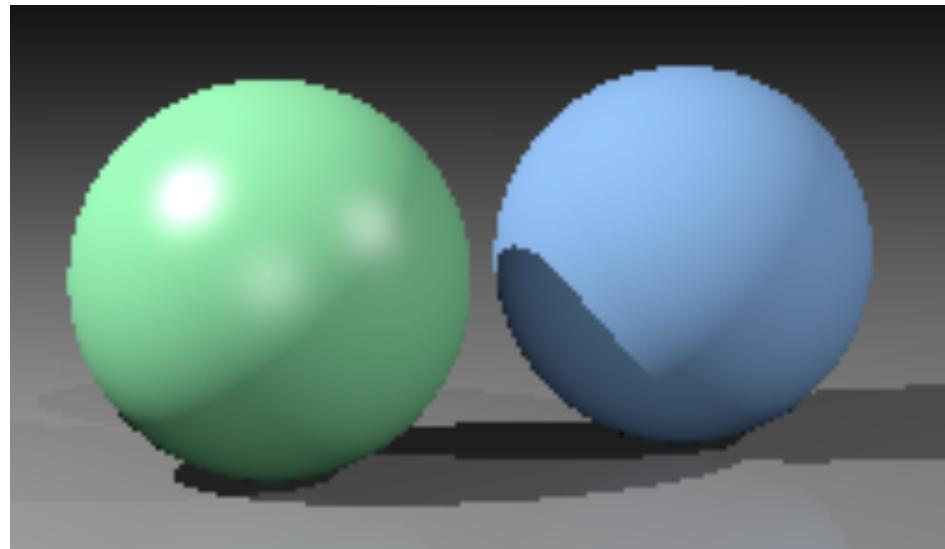
- Antialiasing and compositing both deal with questions of pixels that contain unresolved detail
- Antialiasing: how to carefully throw away the detail
- Compositing: how to account for the detail when combining images

# Aliasing

point sampling a  
continuous image:

continuous image defined  
by ray tracing procedure

continuous image defined  
by a bunch of black rectangles

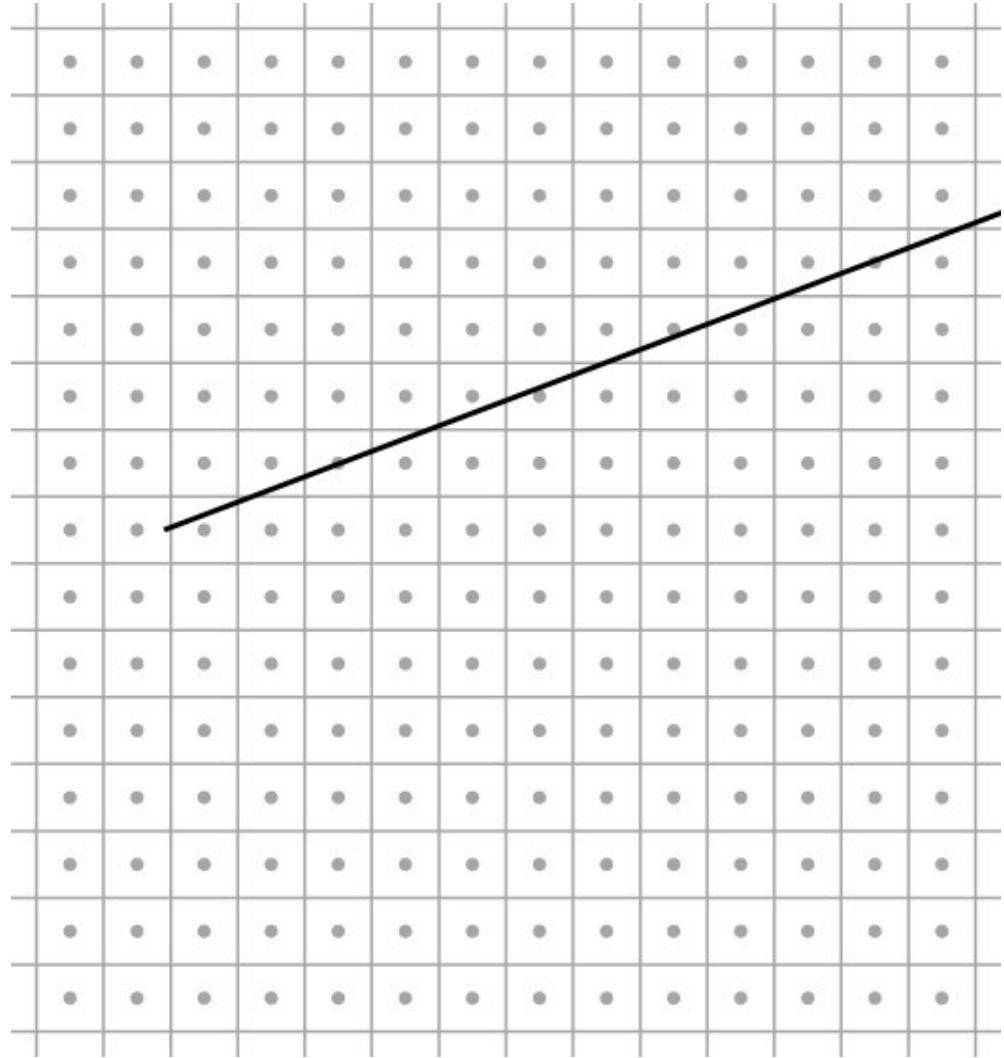


# Antialiasing

- A name for techniques to prevent aliasing
- In image generation, we need to *filter*
  - Boils down to averaging the image over an area
  - Weight by a filter
- Methods depend on source of image
  - Rasterization (lines and polygons)
  - Point sampling (e.g. raytracing)
  - Texture mapping

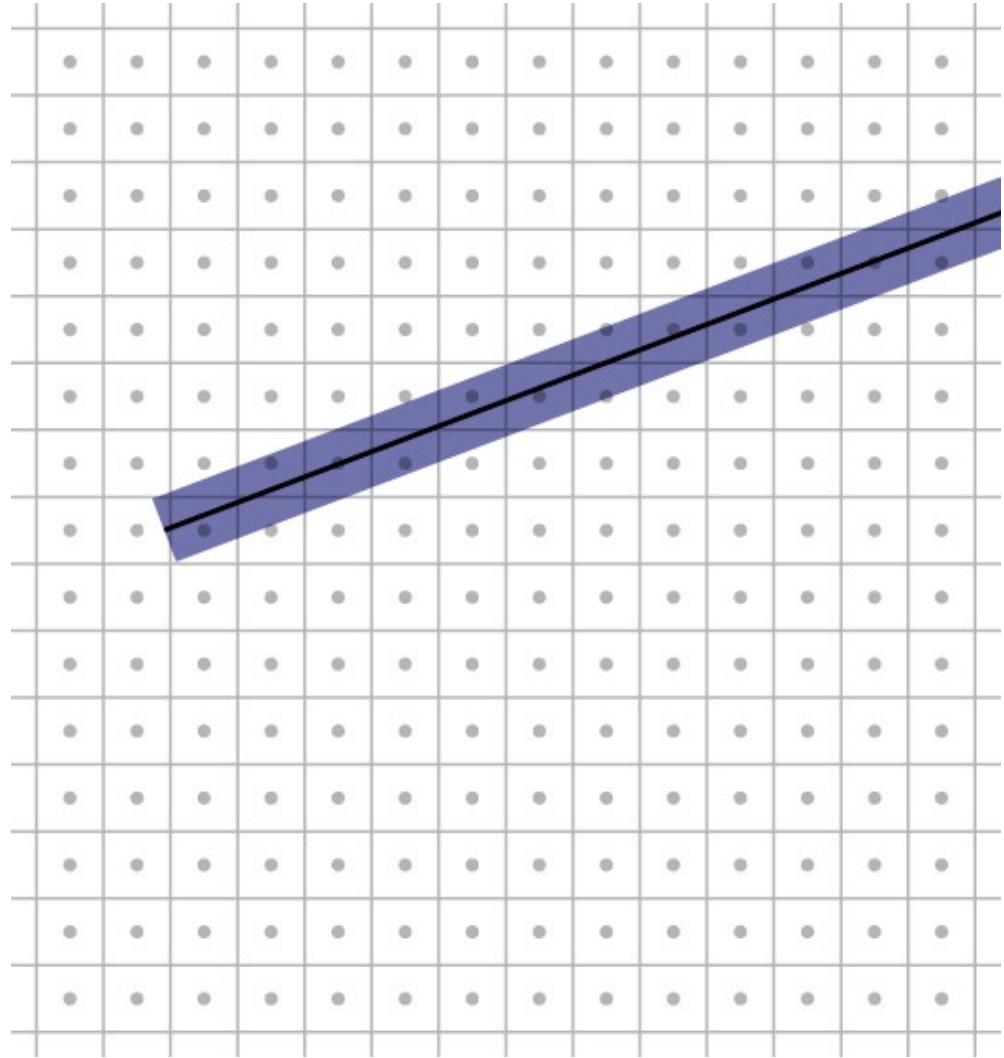
# Rasterizing lines

- Define line as a rectangle
- Specify by two endpoints
- Ideal image: black inside, white outside



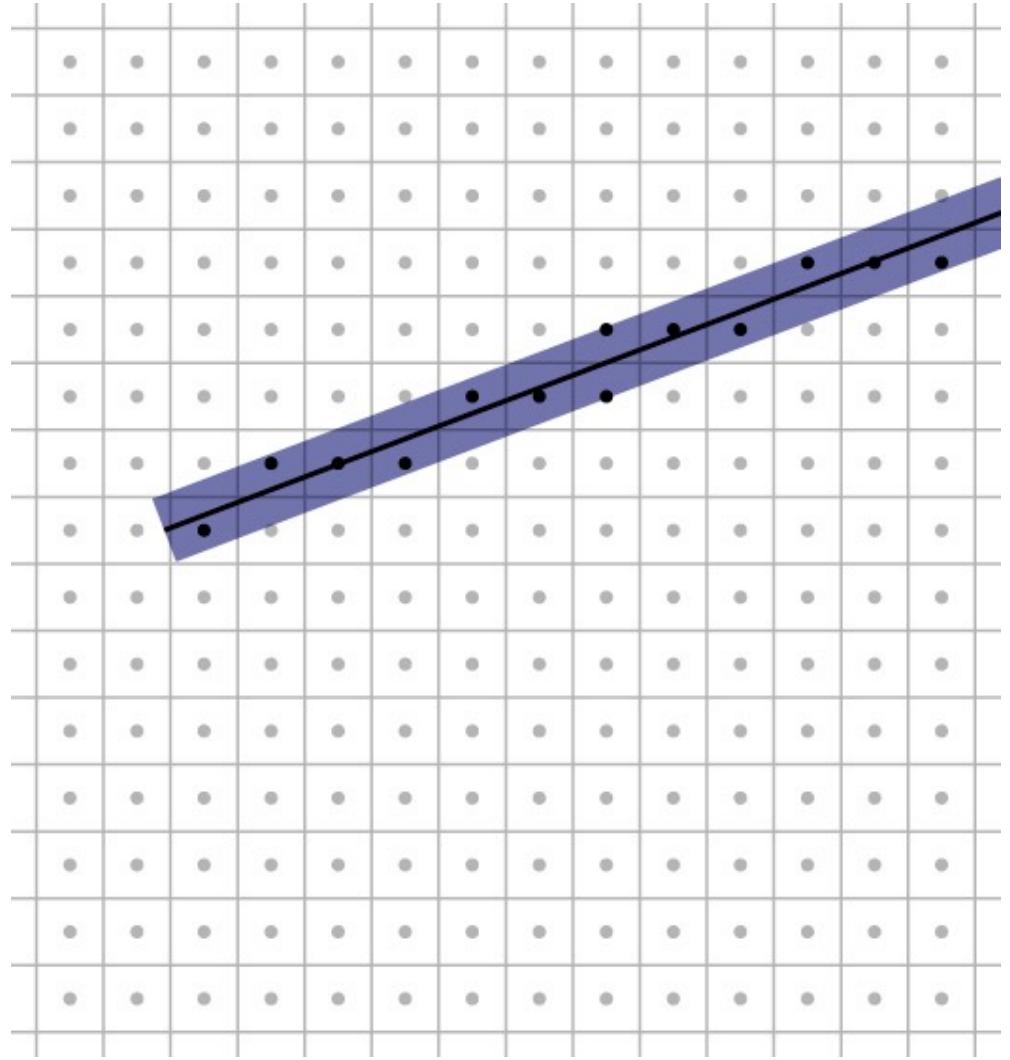
# Rasterizing lines

- Define line as a rectangle
- Specify by two endpoints
- Ideal image: black inside, white outside



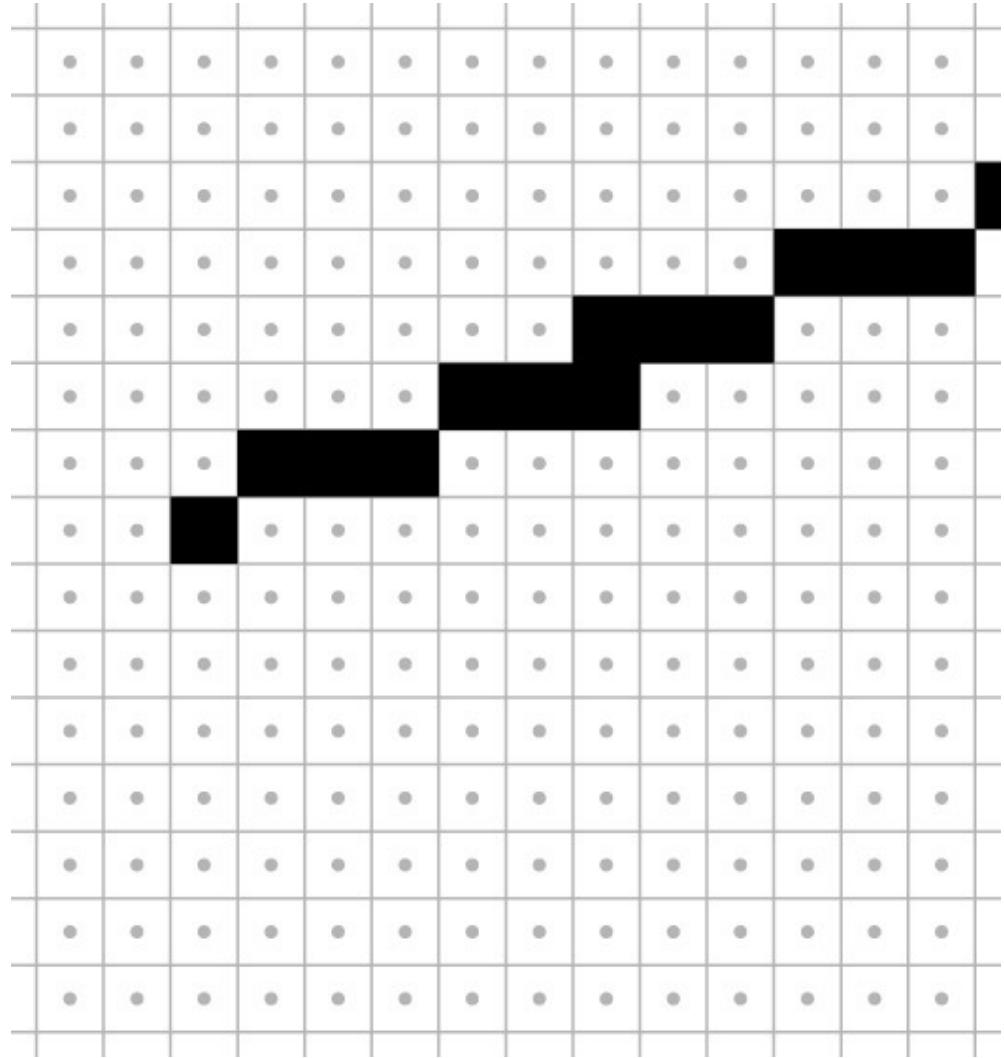
# Point sampling

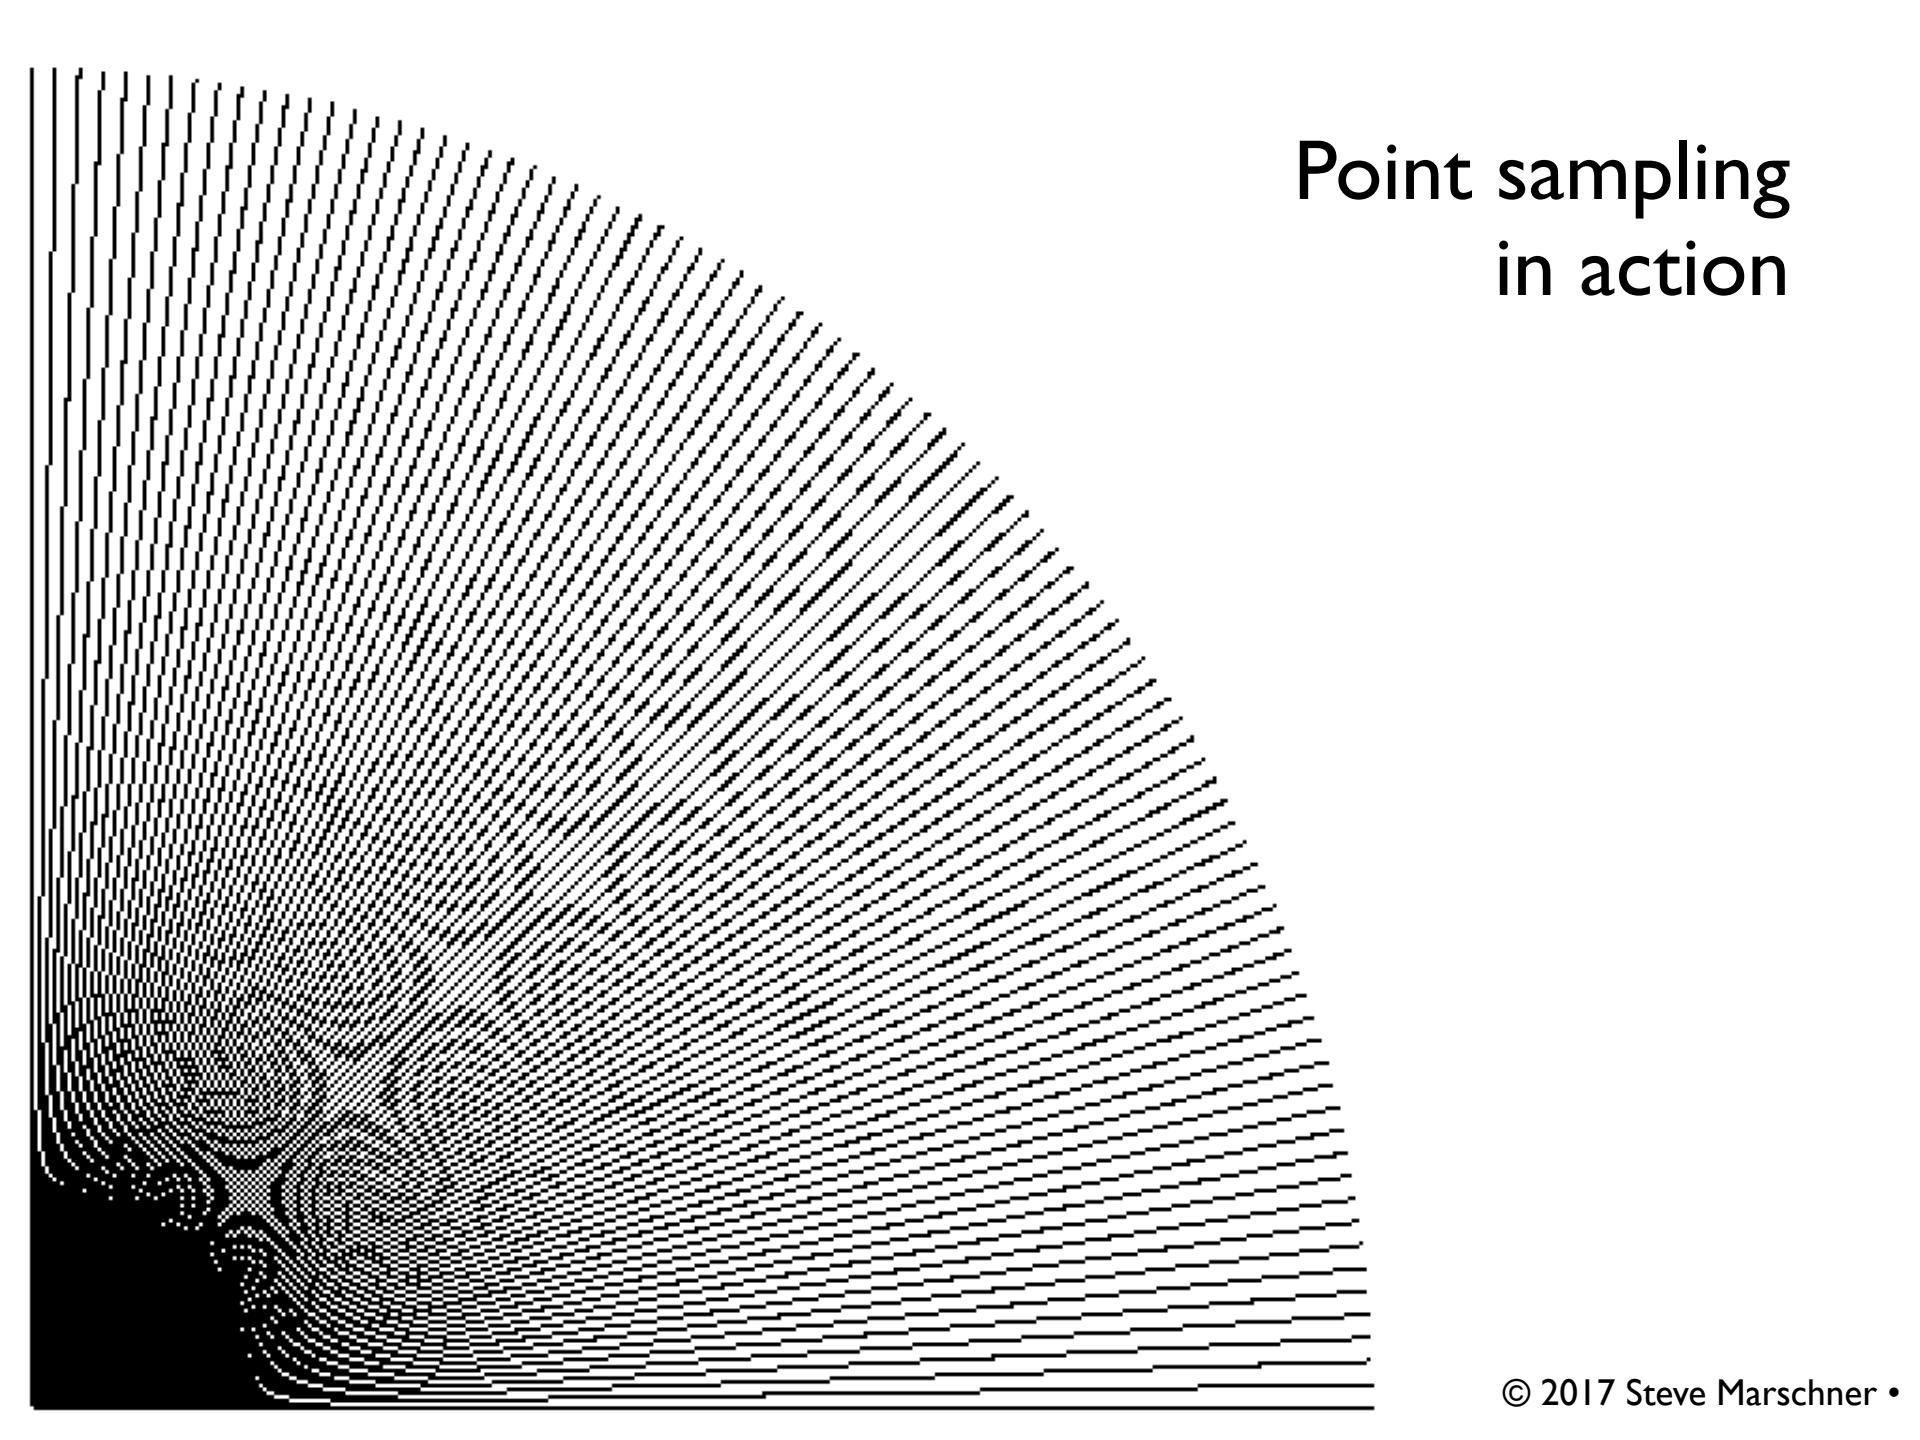
- Approximate rectangle by drawing all pixels whose centers fall within the line
- Problem: all-or-nothing leads to jaggies
  - this is sampling with no filter (aka. point sampling)



# Point sampling

- Approximate rectangle by drawing all pixels whose centers fall within the line
- Problem: all-or-nothing leads to jaggies
  - this is sampling with no filter (aka. point sampling)

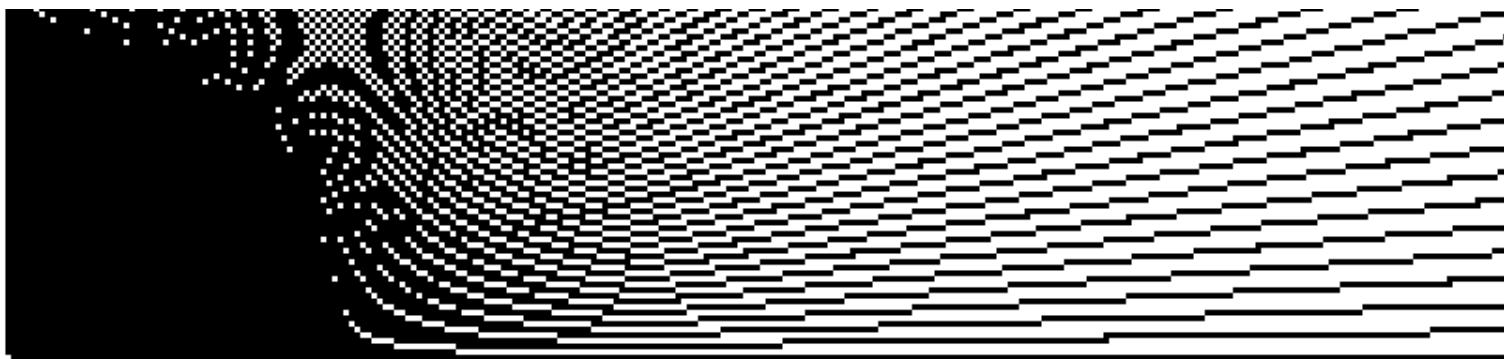




# Point sampling in action

# Aliasing

- Point sampling is fast and simple
- But the lines have stair steps and variations in width
- This is an aliasing phenomenon
  - Sharp edges of line contain high frequencies
- Introduces features to image that are not supposed to be there!

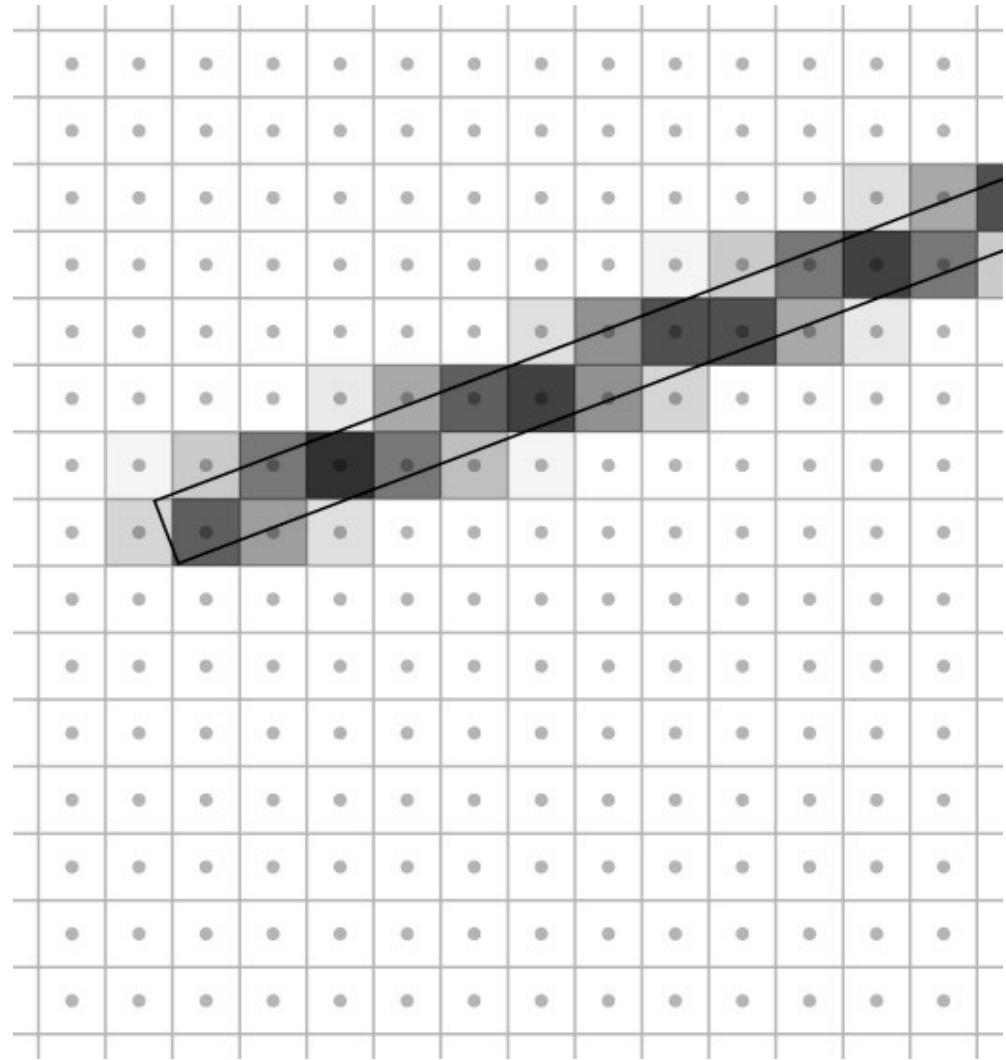


# Antialiasing

- Point sampling makes an all-or-nothing choice in each pixel
  - therefore steps are inevitable when the choice changes
  - yet another example where discontinuities are bad
- On bitmap devices this is necessary
  - hence high resolutions required
  - 600+ dpi in laser printers to make aliasing invisible
- On continuous-tone devices we can do better

# Antialiasing

- Basic idea: replace “is the image black at the pixel center?” with “how much is pixel covered by black?”
- Replace yes/no question with quantitative question.

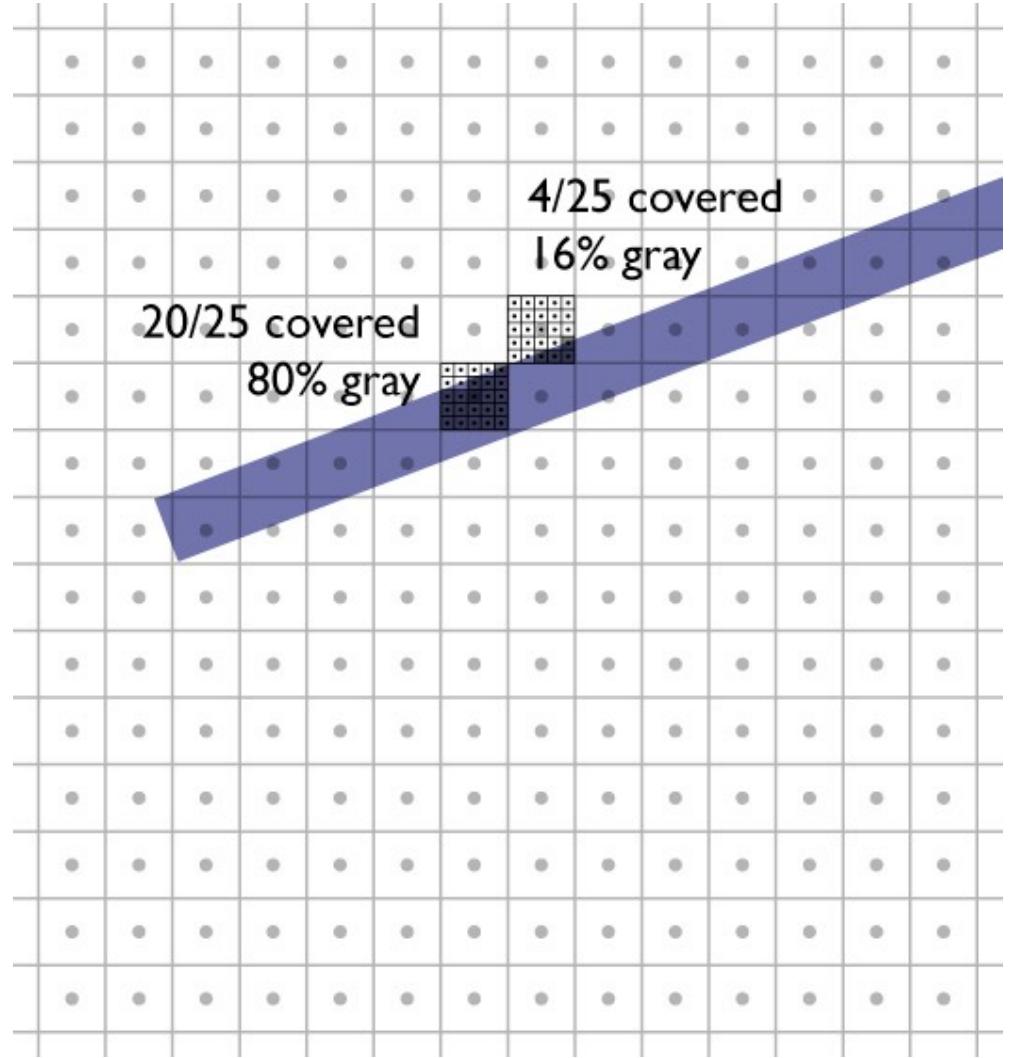


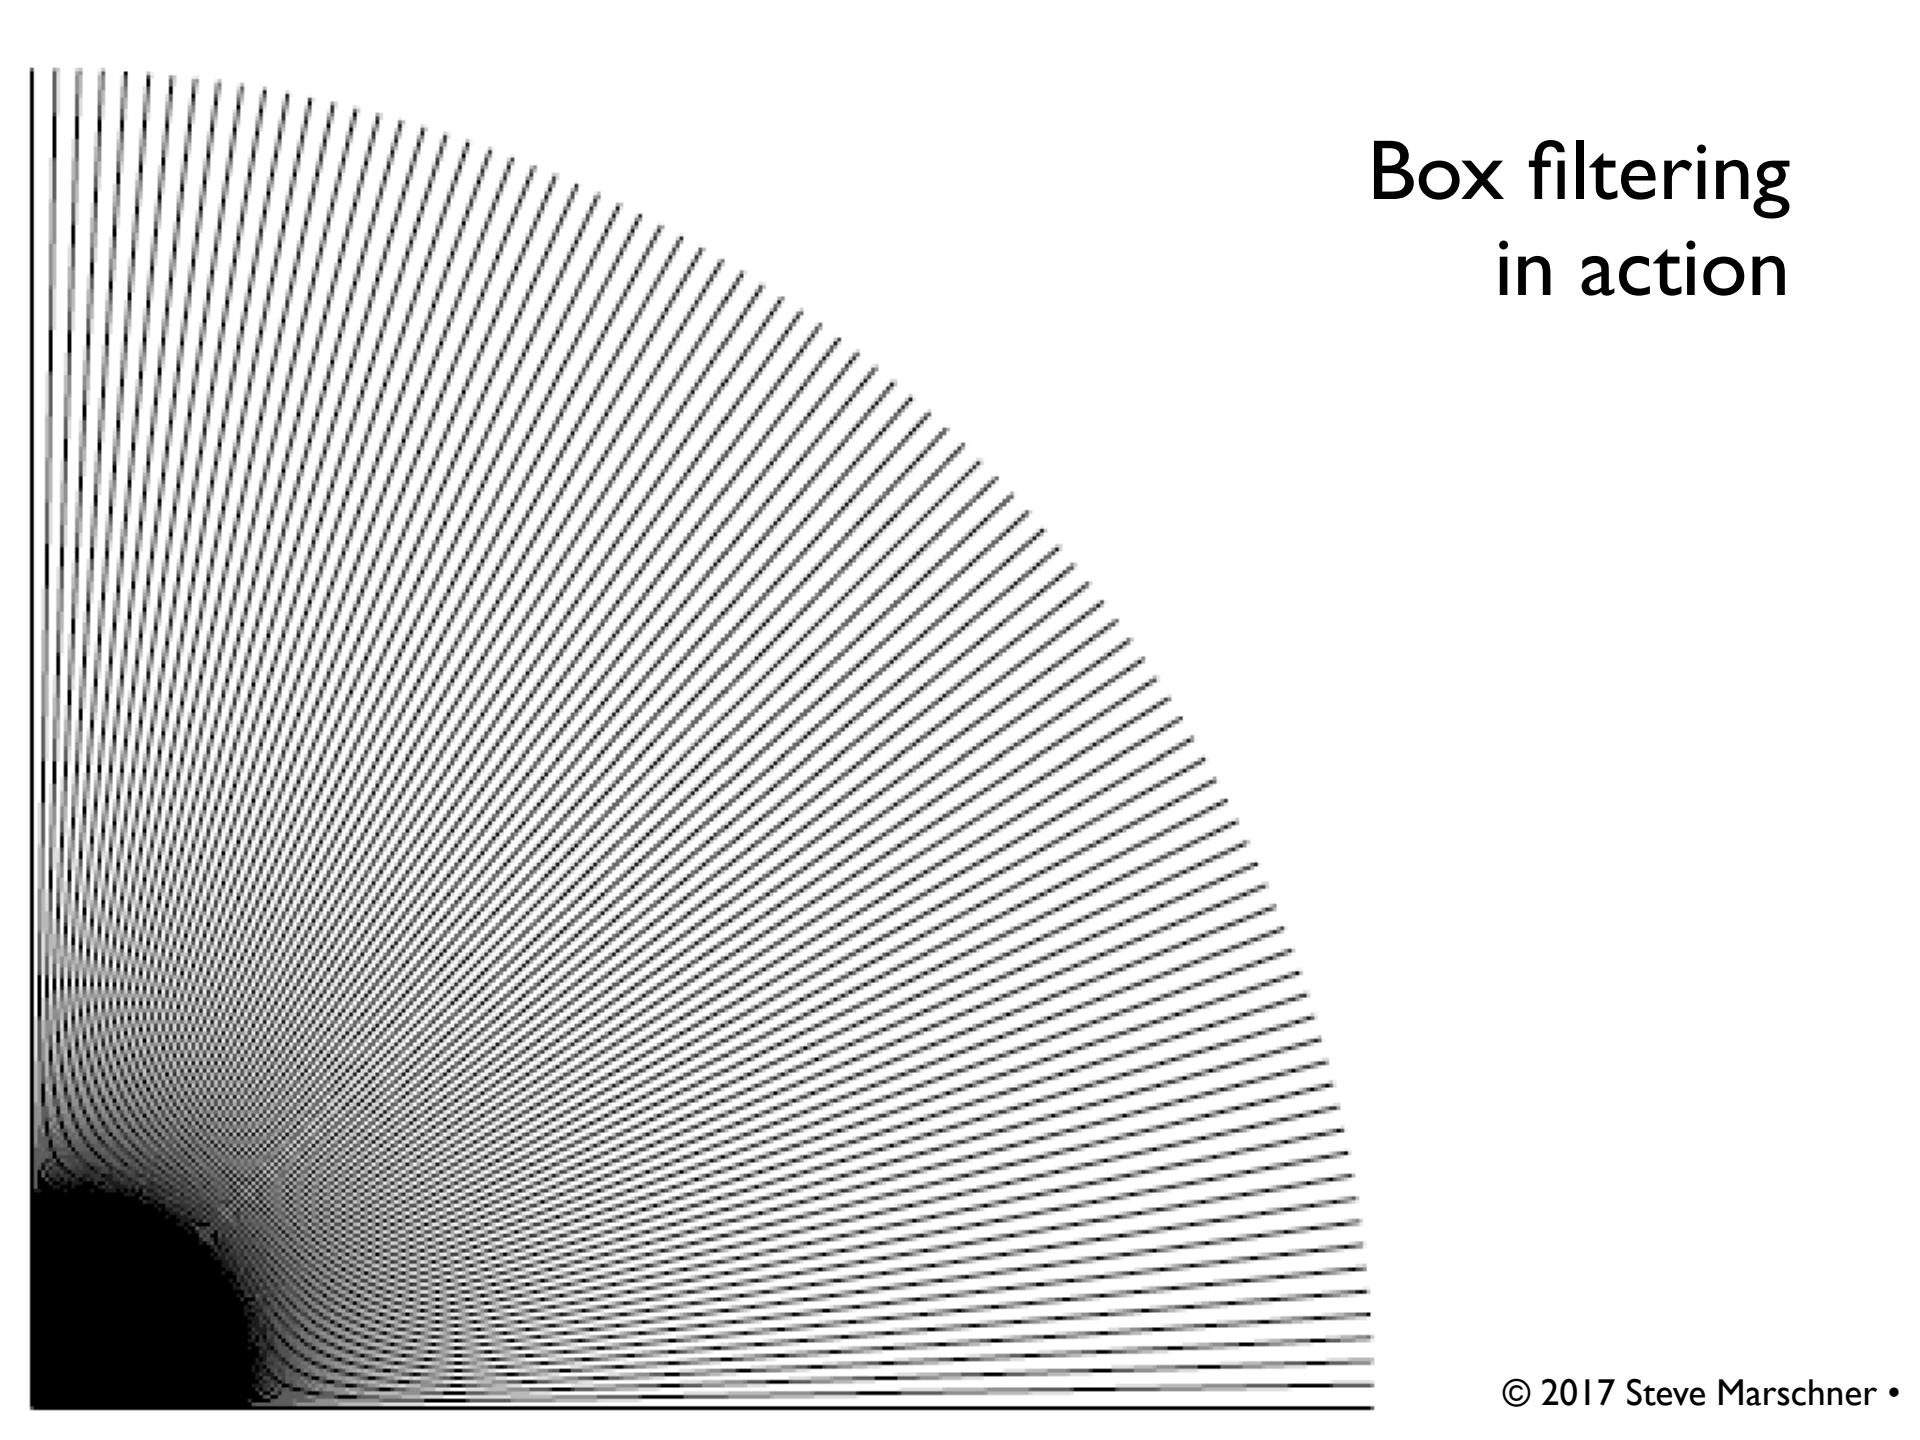
# Box filtering

- Pixel intensity is proportional to area of overlap with square pixel area
- Also called “unweighted area averaging”

# Box filtering by supersampling

- Compute coverage fraction by counting subpixels
- Simple, accurate
- But slow





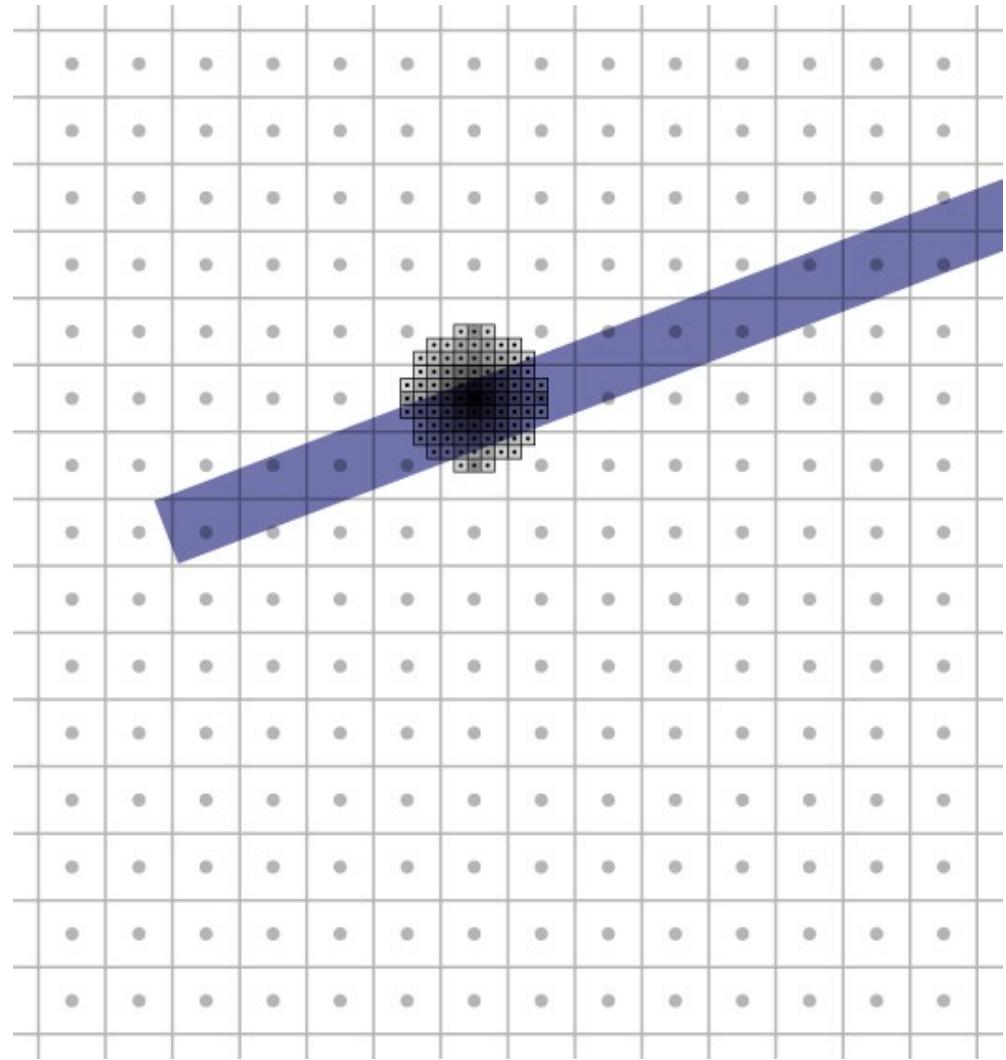
# Box filtering in action

# Weighted filtering

- Box filtering problem: treats area near edge same as area near center
  - results in pixel turning on “too abruptly”
- Alternative: weight area by a smooth function
  - unweighted averaging corresponds to using a box function
  - a gaussian is a popular choice of smooth filter
  - important property: normalization (unit integral)

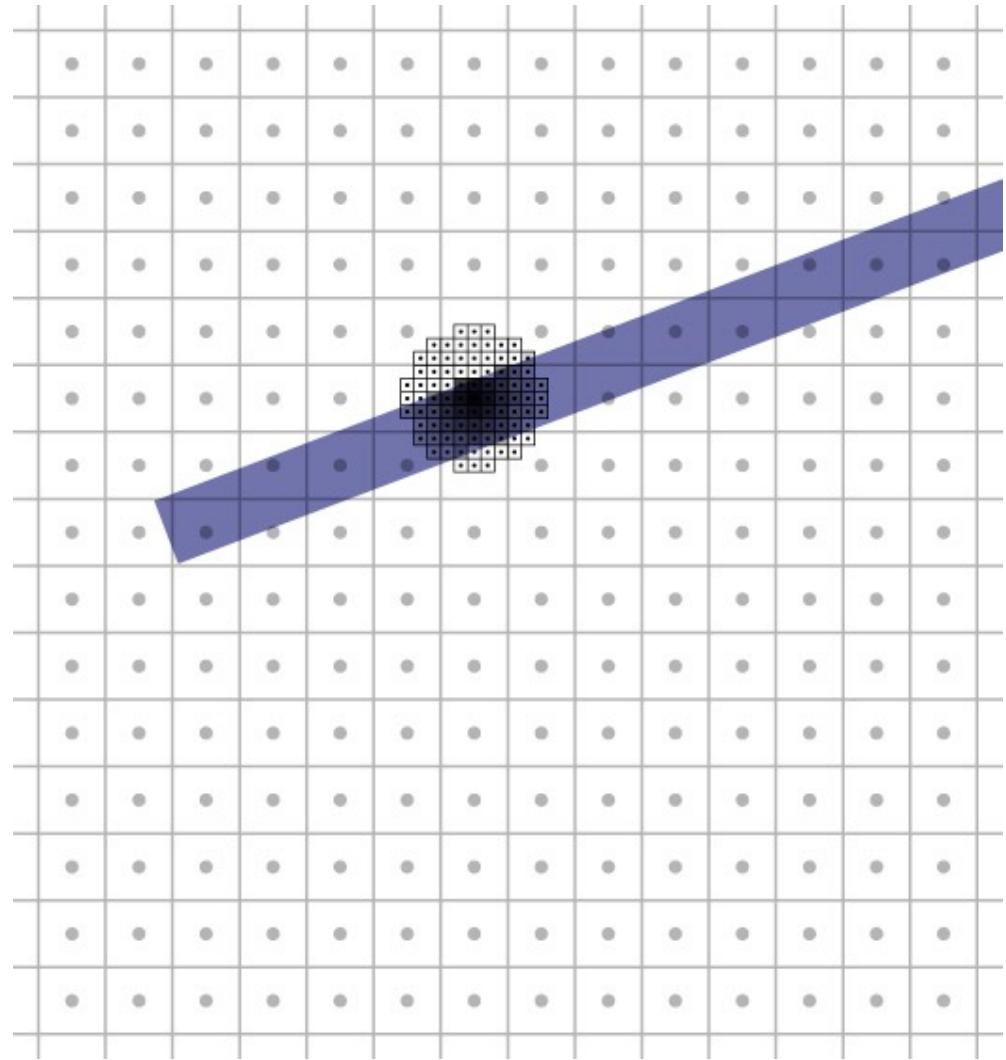
# Weighted filtering by supersampling

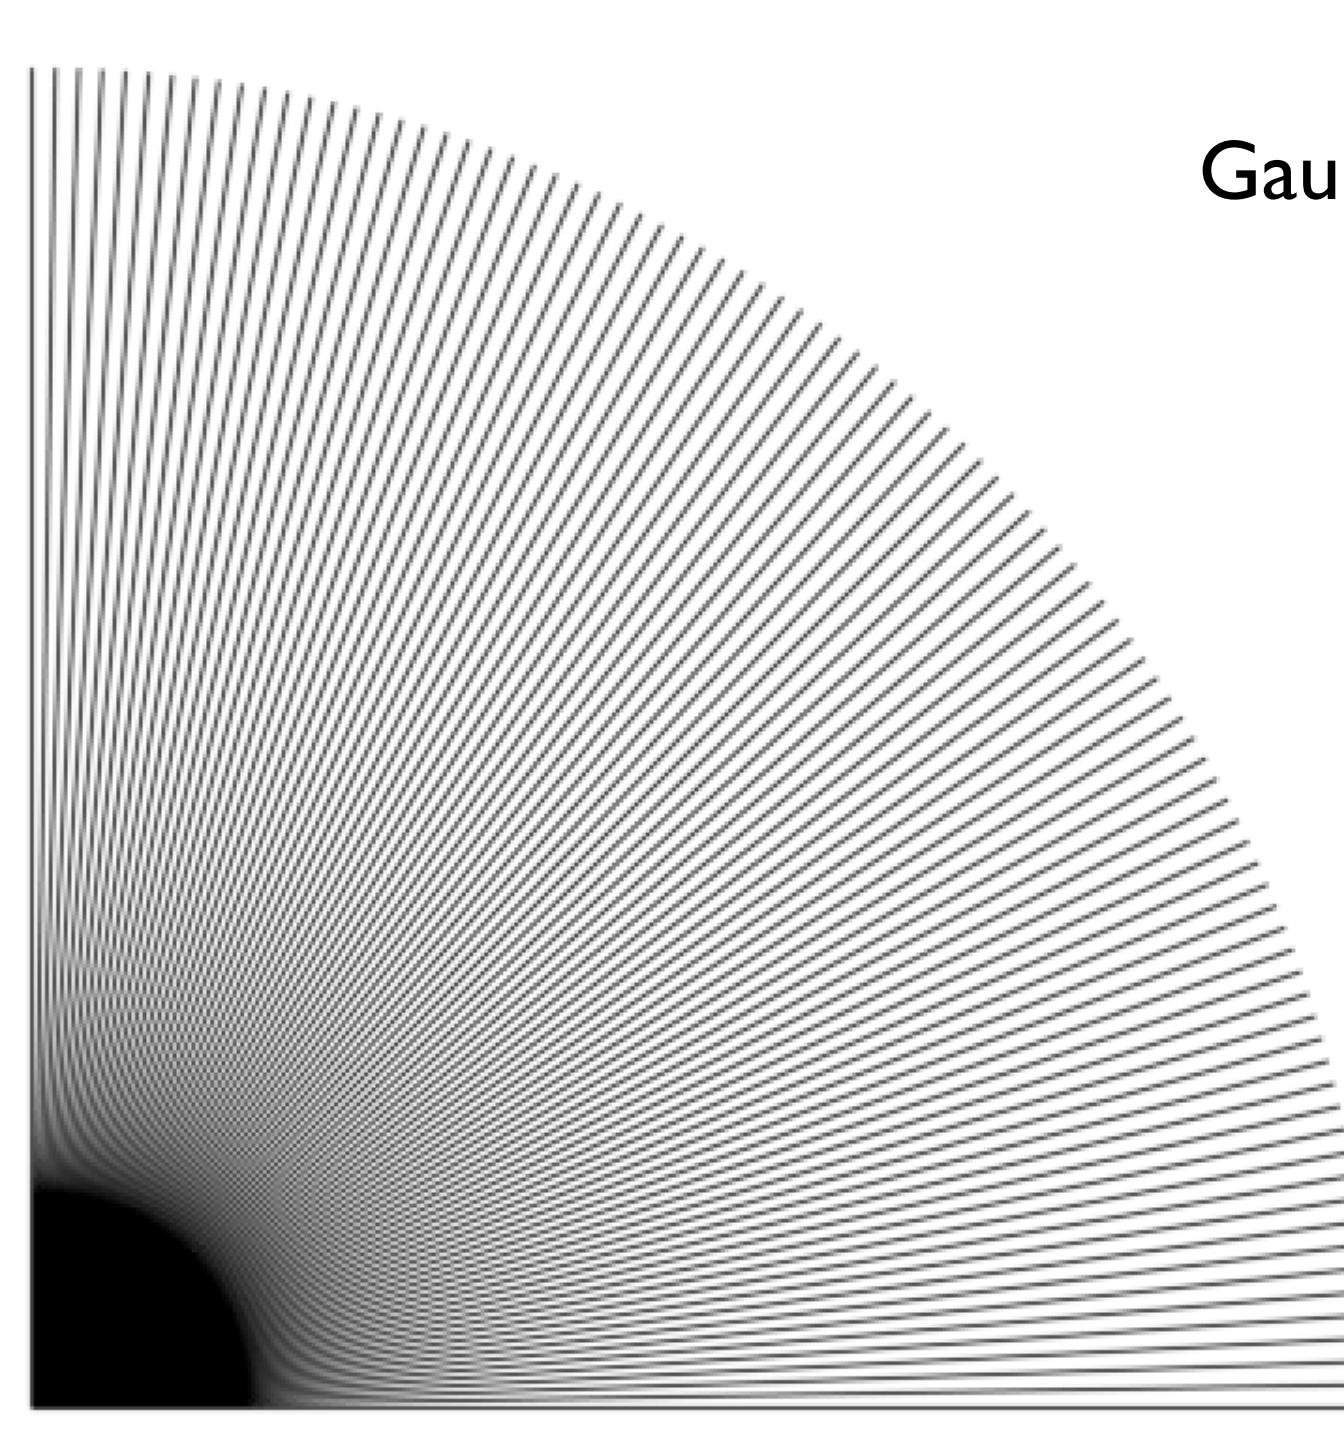
- Compute filtering integral by summing filter values for covered subpixels
- Simple, accurate
- But really slow



# Weighted filtering by supersampling

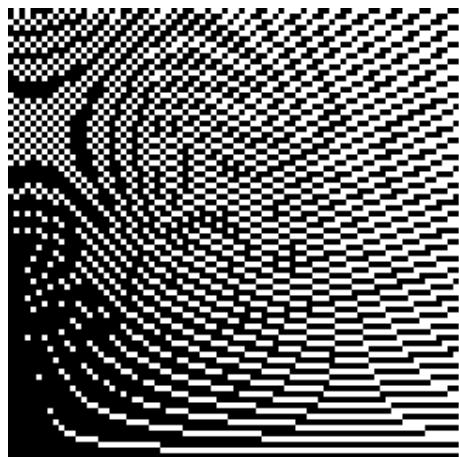
- Compute filtering integral by summing filter values for covered subpixels
- Simple, accurate
- But really slow



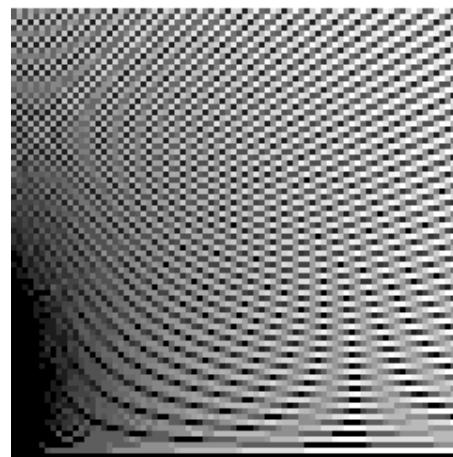


# Gaussian filtering in action

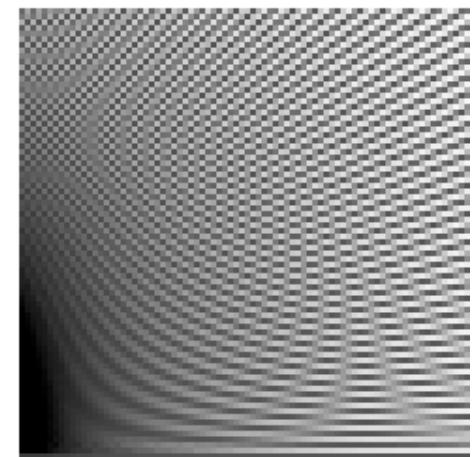
# Filter comparison



Point sampling

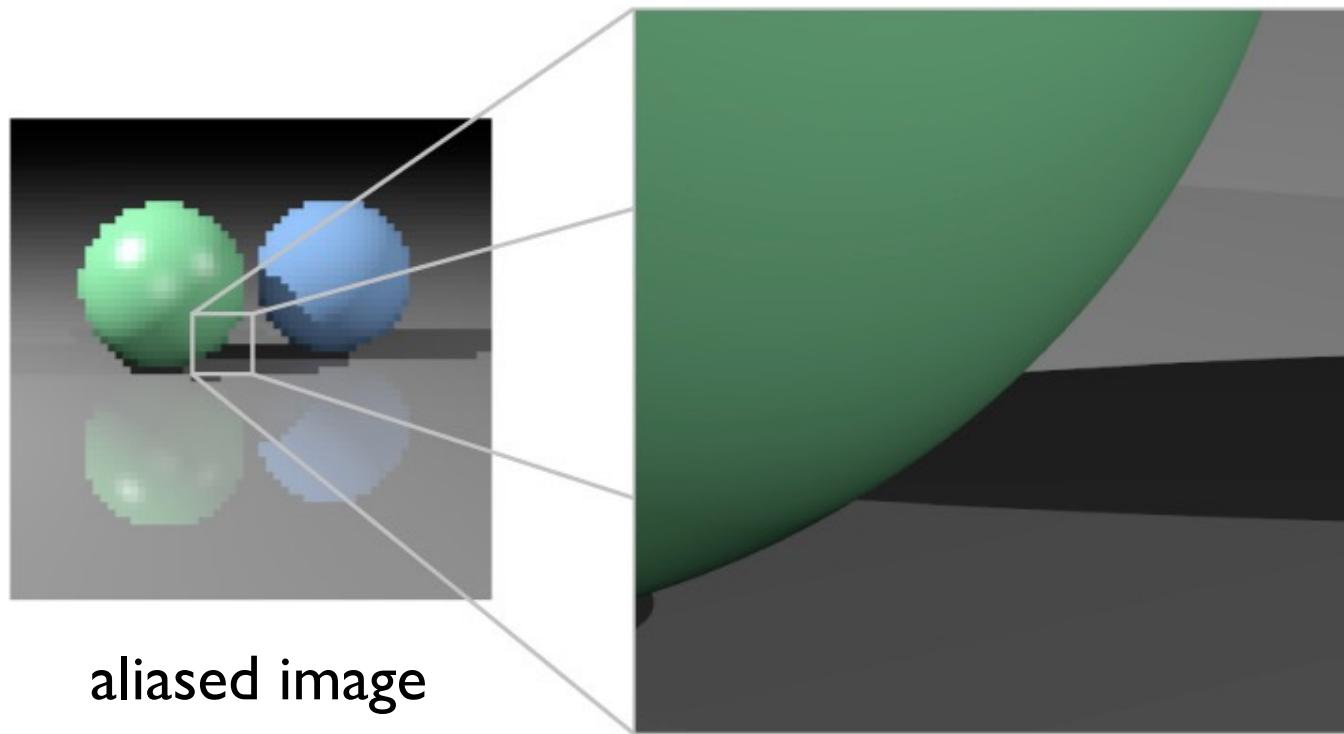


Box filtering

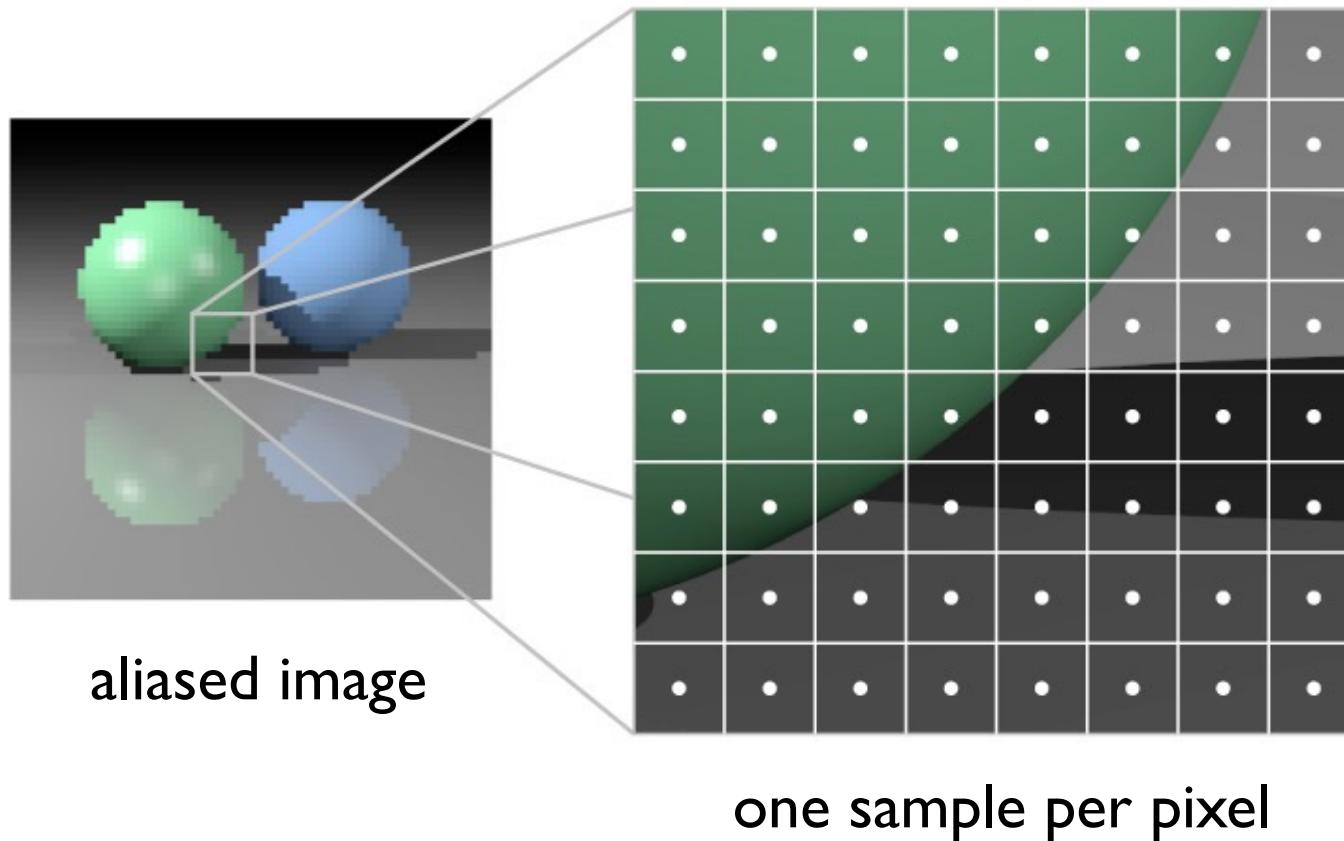


Gaussian filtering

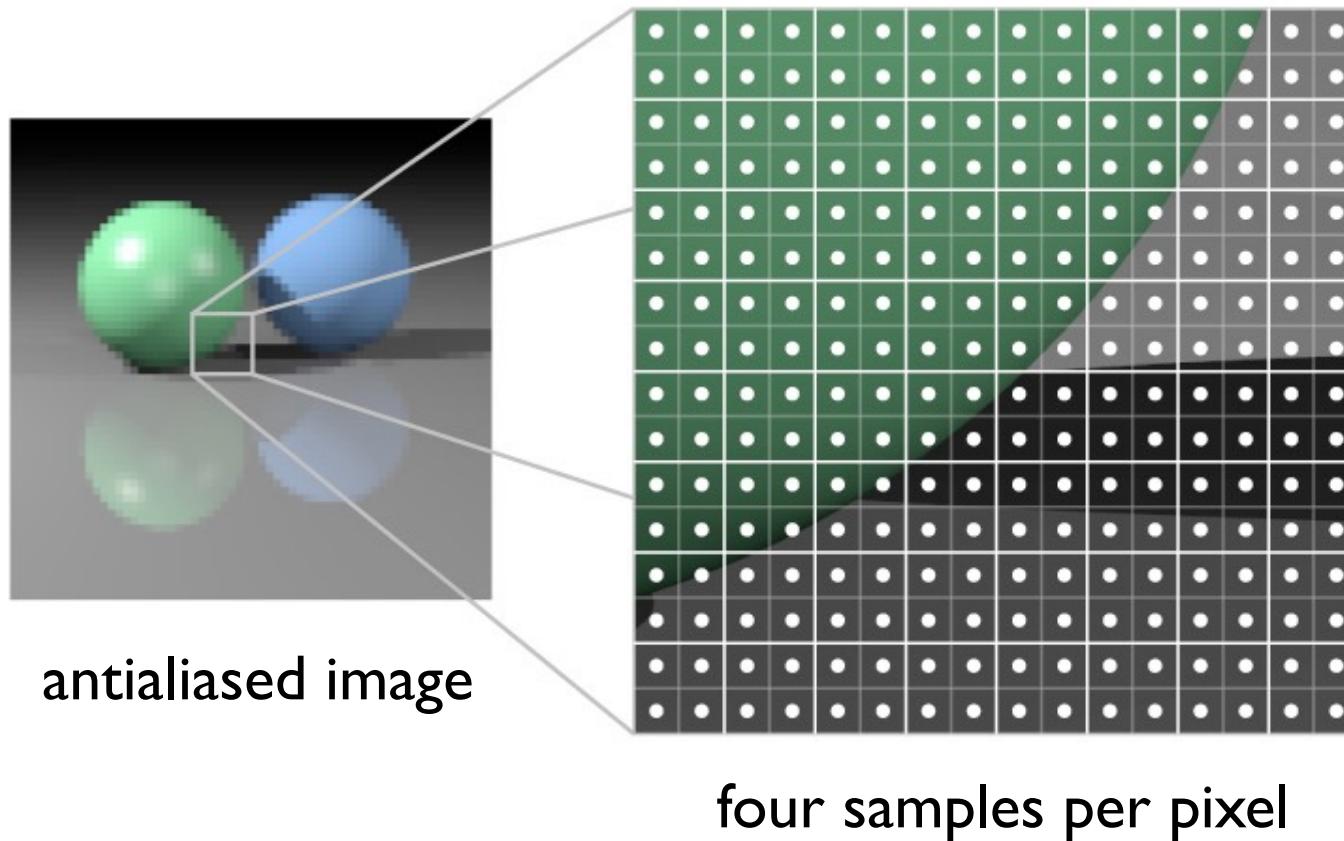
# Antialiasing in ray tracing



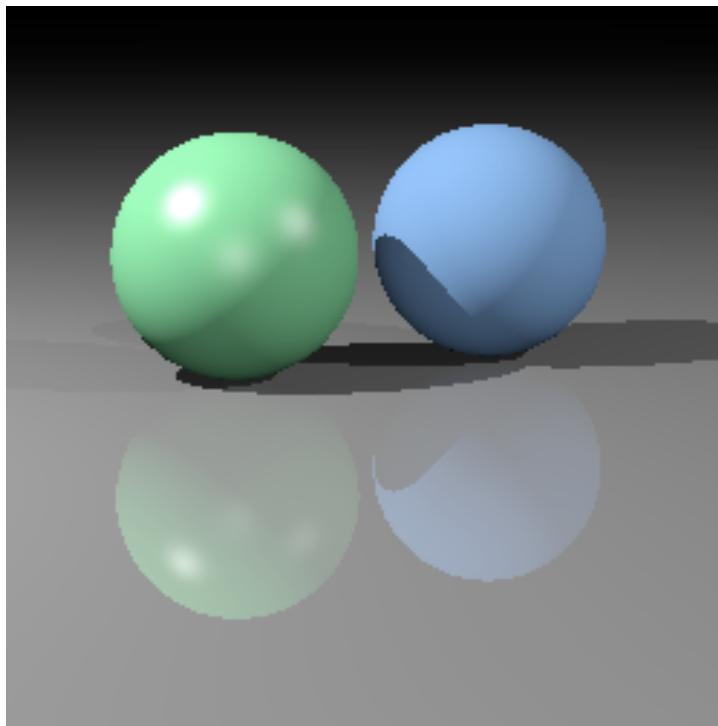
# Antialiasing in ray tracing



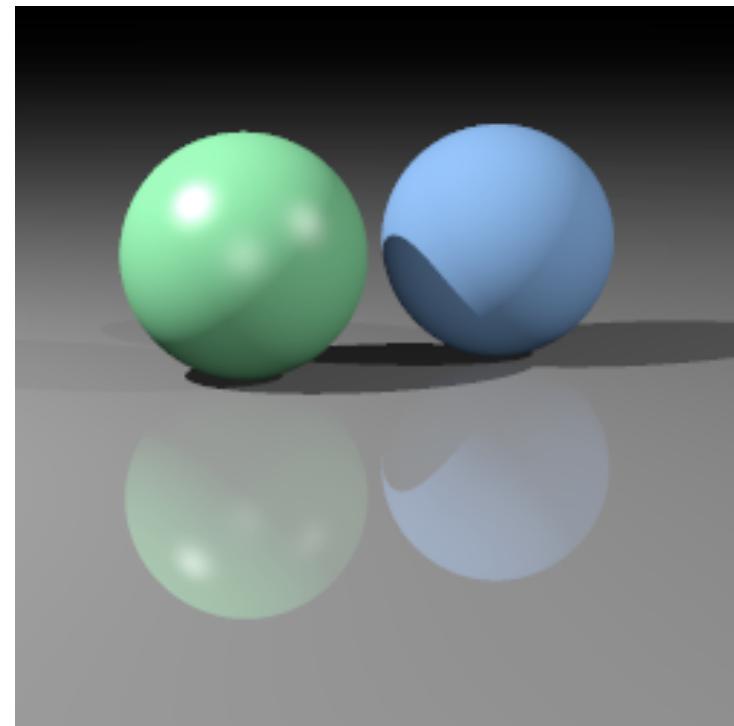
# Antialiasing in ray tracing



# Antialiasing in ray tracing



one sample/pixel

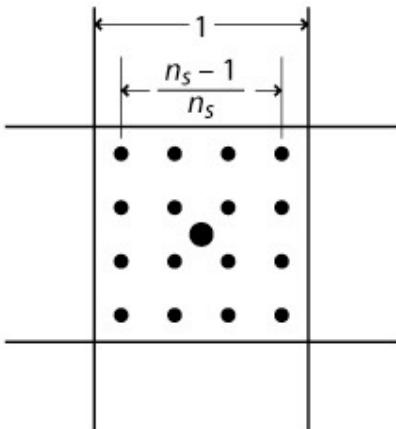


9 samples/pixel

# Details of supersampling

- For image coordinates with integer pixel centers:

```
// one sample per pixel  
for iy = 0 to (ny-1) by 1  
  for ix = 0 to (nx-1) by 1 {  
    ray = camera.getRay(ix, iy);  
    image.set(ix, iy, trace(ray));  
  }
```



```
// ns^2 samples per pixel  
for iy = 0 to (ny-1) by 1  
  for ix = 0 to (nx-1) by 1 {  
    Color sum = 0;  
    for dx = -(ns-1)/2 to (ns-1)/2 by 1  
      for dy = -(ns-1)/2 to (ns-1)/2 by 1 {  
        x = ix + dx / ns;  
        y = iy + dy / ns;  
        ray = camera.getRay(x, y);  
        sum += trace(ray);  
      }  
    image.set(ix, iy, sum / (ns*ns));  
  }
```

# Details of supersampling

- For image coordinates in unit square

```
// one sample per pixel
for iy = 0 to (ny-1) by 1
    for ix = 0 to (nx-1) by 1 {
        double x = (ix + 0.5) / nx;
        double y = (iy + 0.5) / ny;
        ray = camera.getRay(x, y);
        image.set(ix, iy, trace(ray));
    }
```

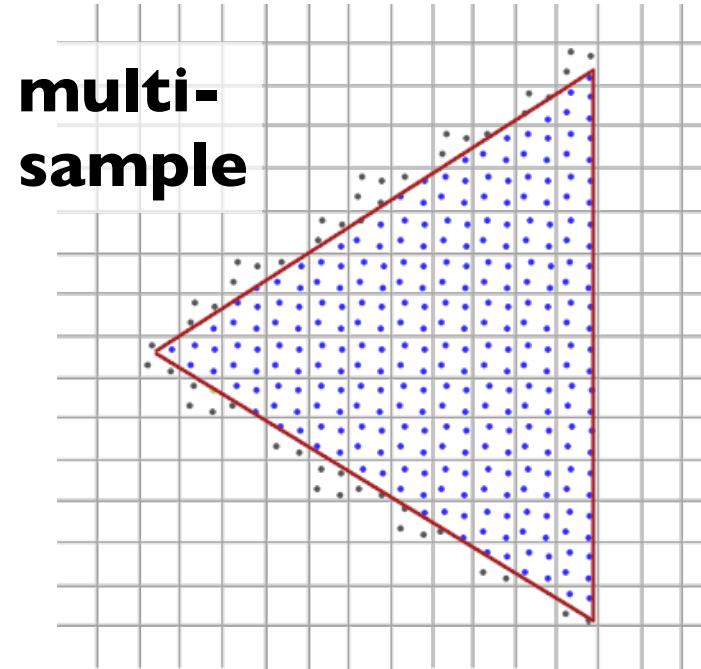
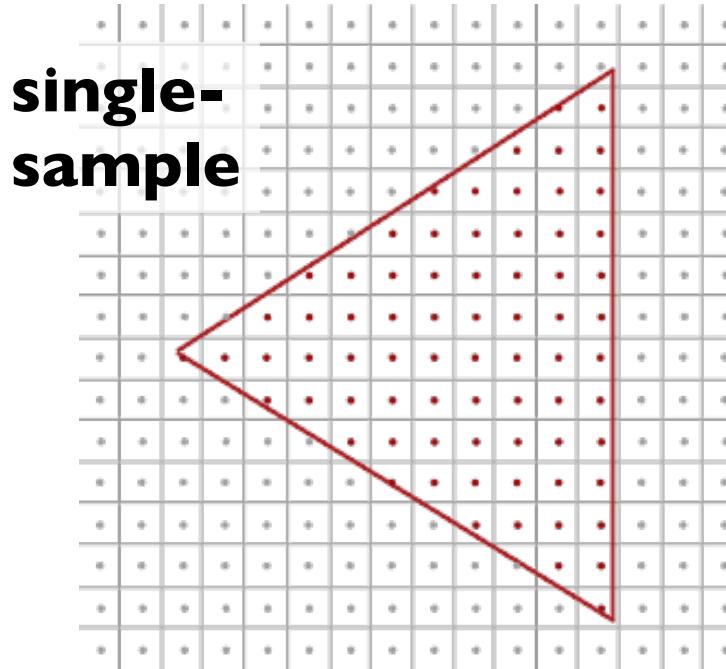
```
// ns^2 samples per pixel
for iy = 0 to (ny-1) by 1
    for ix = 0 to (nx-1) by 1 {
        Color sum = 0;
        for dx = 0 to (ns-1) by 1
            for dy = 0 to (ns-1) by 1 {
                x = (ix + (dx + 0.5) / ns) / nx;
                y = (iy + (dy + 0.5) / ns) / ny;
                ray = camera.getRay(x, y);
                sum += trace(ray);
            }
        image.set(ix, iy, sum / (ns*ns));
    }
```

# Supersampling vs. multisampling

- Supersampling is terribly expensive
- GPUs use an approximation called *multisampling*
  - Compute one shading value per pixel
  - Store it at many subpixel samples, each with its own depth

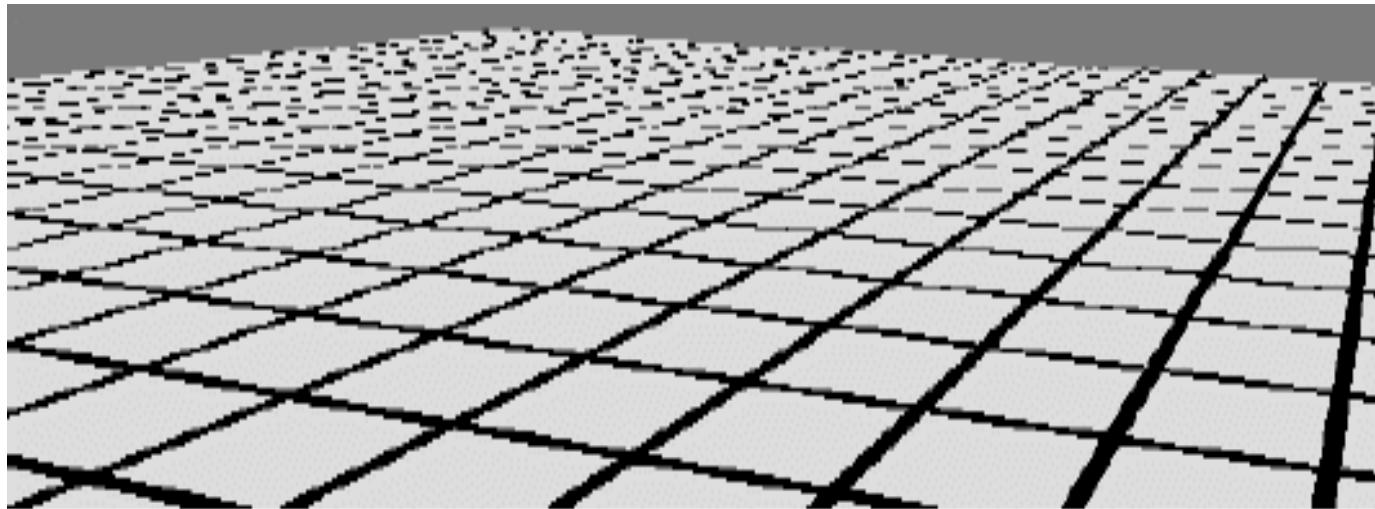
# Multisample rasterization

- Each fragment carries several (color,depth) samples
  - shading is computed per-fragment
  - depth test is resolved per-sample
  - final color is average of sample colors



# Antialiasing in textures

- Even with multisampling, we still only evaluate textures once per fragment
- Need to filter the texture somehow!
  - perspective produces very fine subpixel detail



[Akenine-Möller & Haines 2002]

# Texture mapping from 0 to infinity

- When you go close...

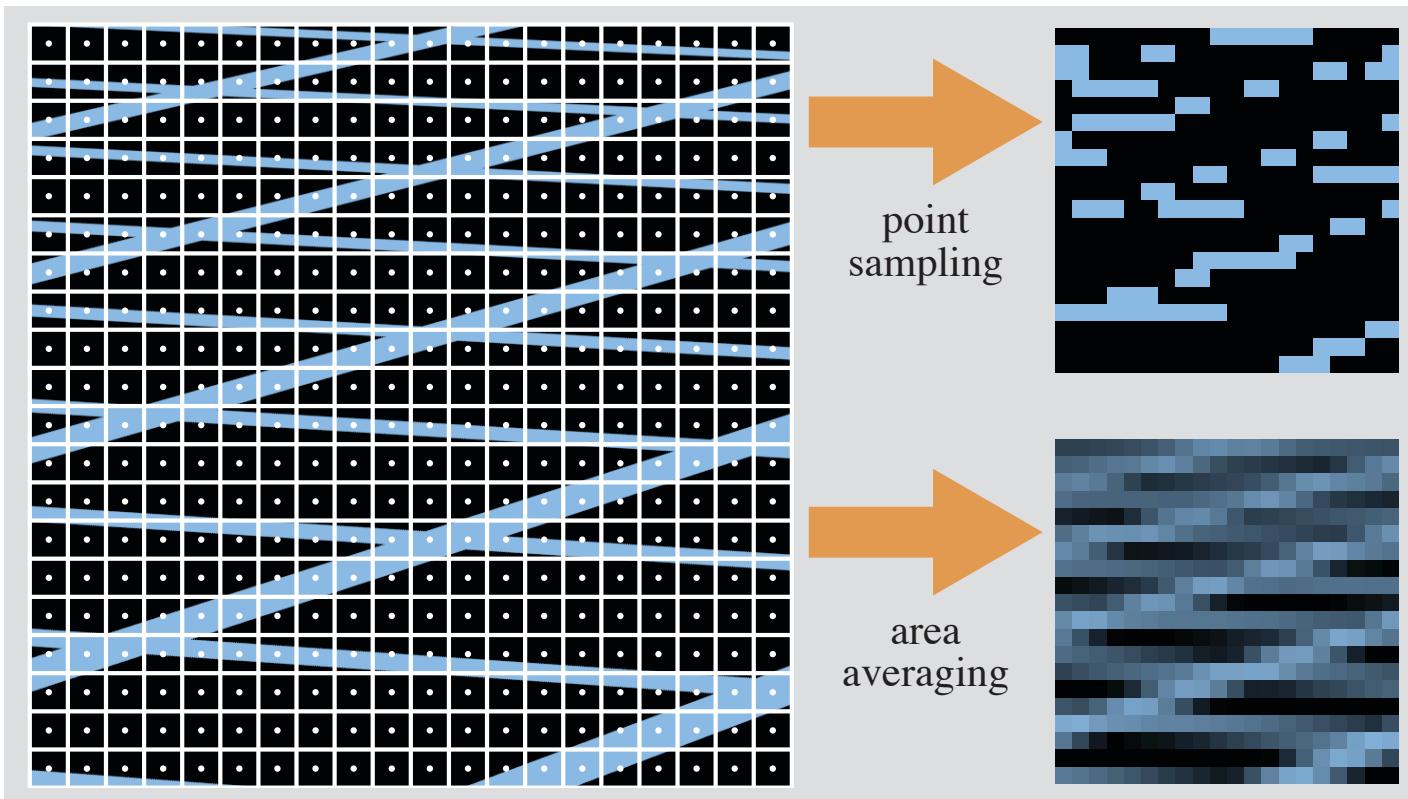


# Texture mapping from 0 to infinity

- When you go far...



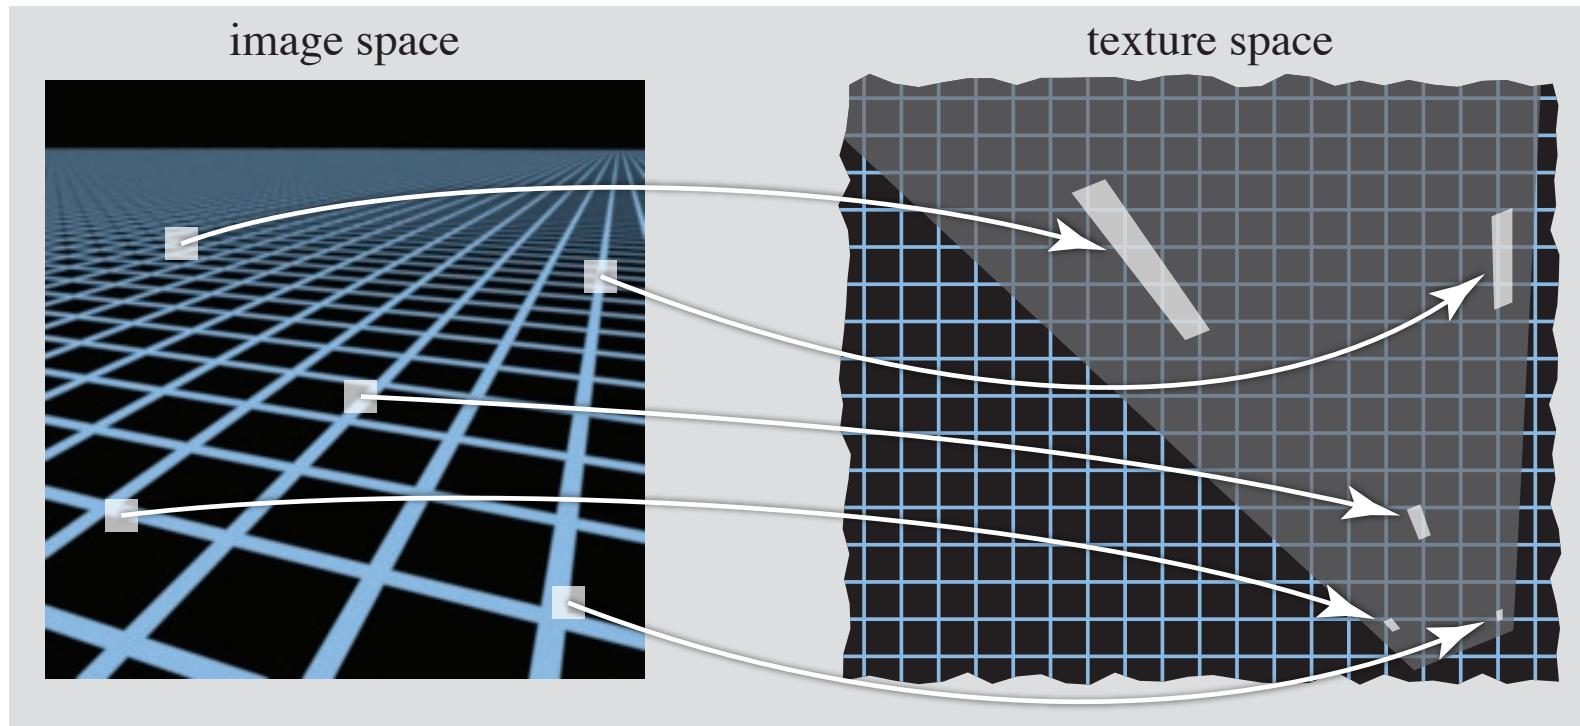
# Solution: pixel filtering



# Pixel filtering in texture space

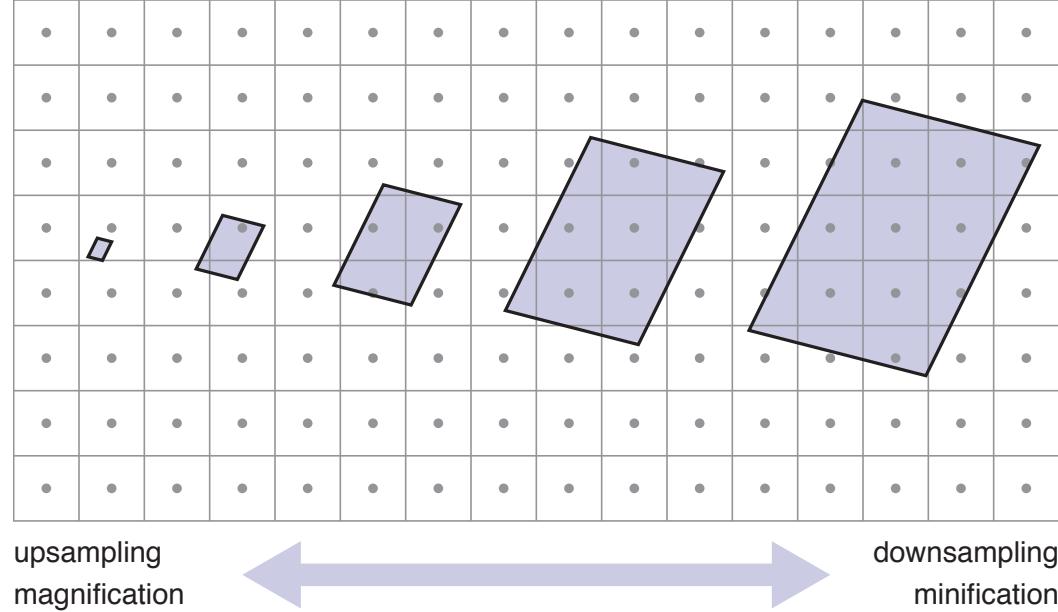
- Sampling is happening in image space
  - sample is a weighted average over a pixel-sized area
  - uniform, predictable, friendly problem!
- Signal is defined in texture space
  - mapping between image and texture is nonuniform
  - each sample is a weighted average over a different sized and shaped area
  - irregular, unpredictable, unfriendly!
- This is a change of variable
  - integrate over texture coordinates rather than image coordinates

# Pixel footprints

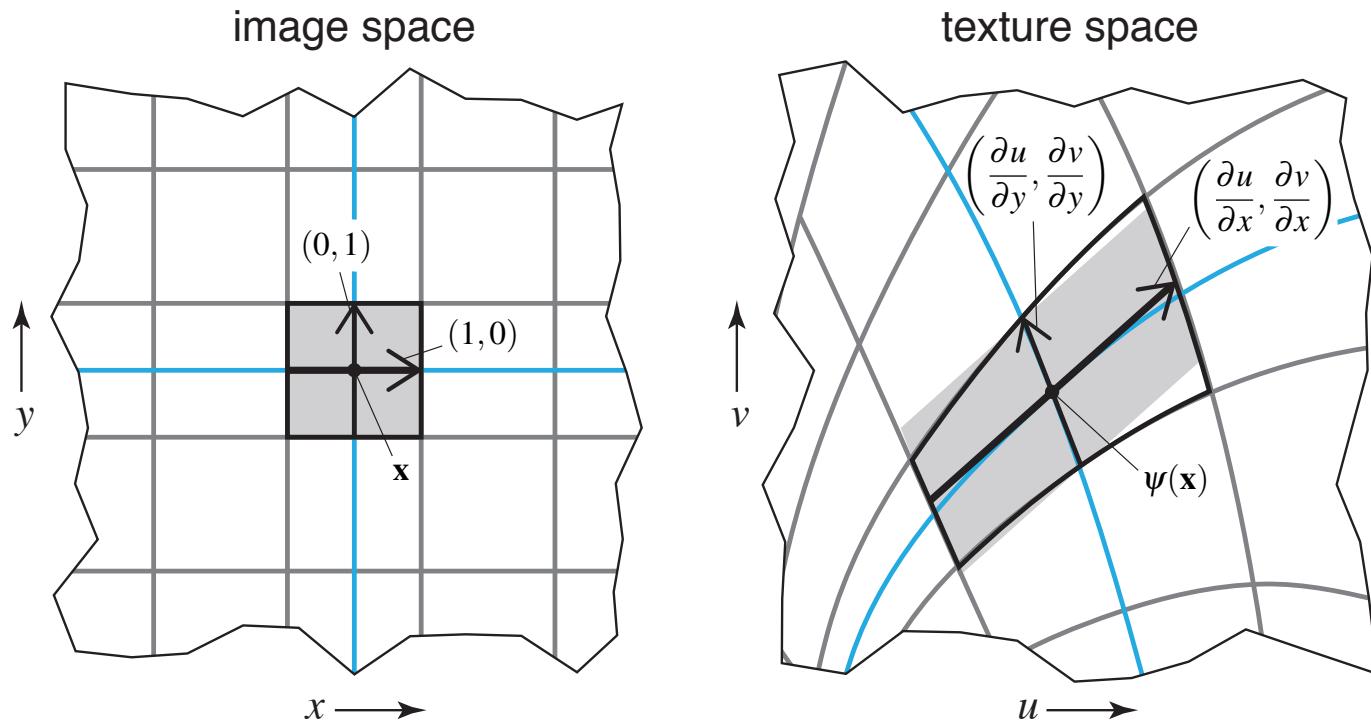


# How does area map over distance?

- At optimal viewing distance:
  - One-to-one mapping between pixel area and texel area
- When closer
  - Each pixel is a small part of the texel
  - magnification
- When farther
  - Each pixel could include many texels
  - “minification”

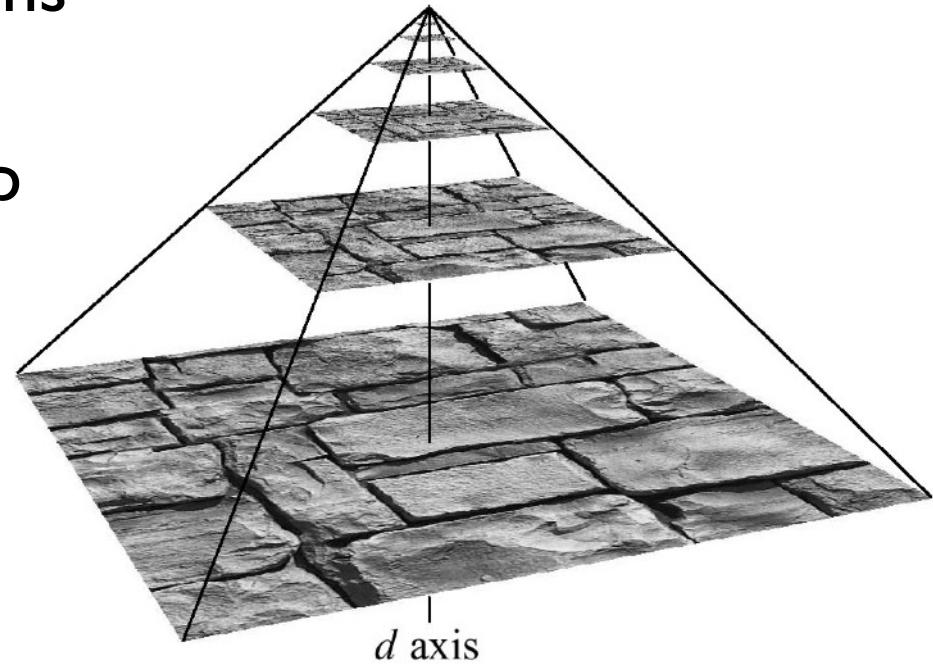


# Sizing up the situation with the Jacobian



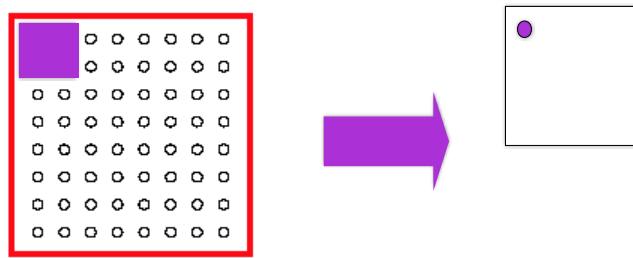
# Mipmap image pyramid

- MIP Maps
  - Multum in Parvo: Much in little, many in small places
  - Proposed by Lance Williams
- Stores pre-filtered versions of texture
- Supports very fast lookup
  - but only of circular filters at certain scales



[Akenine-Möller & Haines 2002]

# Filtering by Averaging



- Each pixel in a level corresponds to 4 pixels in lower level
  - Average
  - Gaussian filtering

# Using the MIP Map

- Find the MIP Map level where the pixel has a 1-to-1 mapping
- How?
  - Find largest side of pixel footprint in texture space
    - Pick level where that side corresponds to a texel
  - Compute derivatives to find pixel footprint

- x derivative:

$$\frac{\partial u}{\partial x} \quad \frac{\partial v}{\partial x}$$

- y derivative:

$$\frac{\partial u}{\partial y} \quad \frac{\partial v}{\partial y}$$

# Given derivatives: what is level?

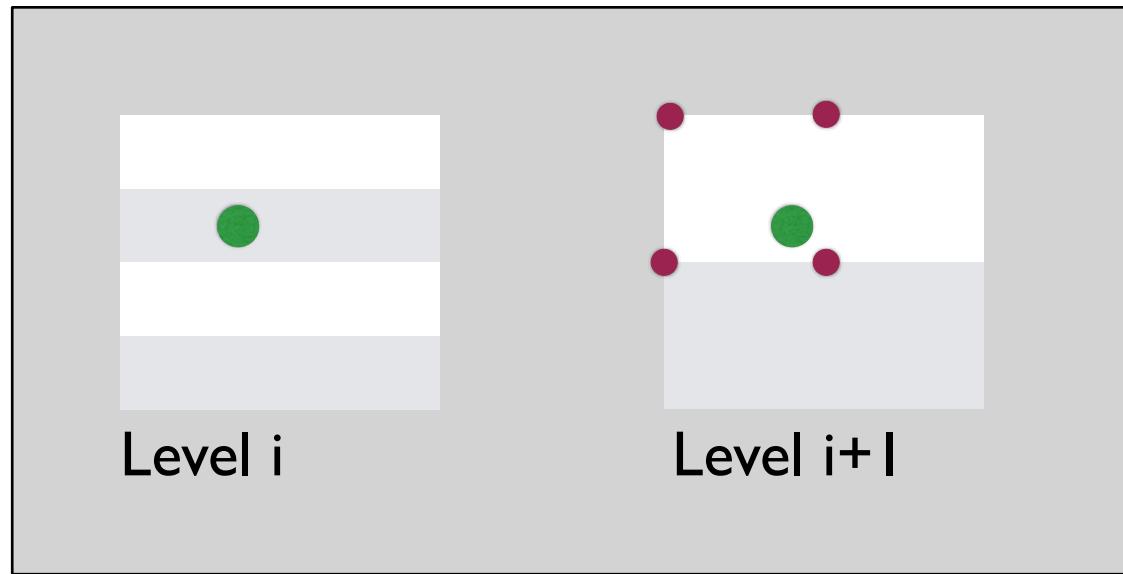
- Derivatives
  - Available in pixel shader (except where there is dynamic branching)

$$level = \log[\max\left(\frac{du}{dx}, \frac{dv}{dx}, \frac{du}{dy}, \frac{dv}{dy}\right)]$$

$$level = \log \sqrt{\left(\frac{du}{dx}\right)^2 + \left(\frac{dv}{dx}\right)^2 + \left(\frac{du}{dy}\right)^2 + \left(\frac{dv}{dy}\right)^2}$$

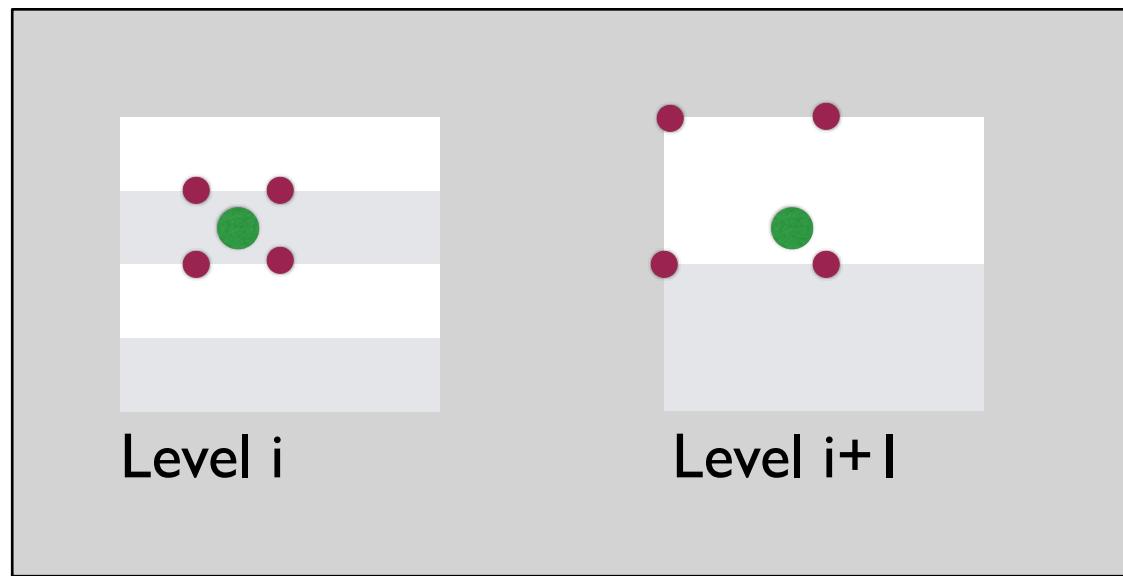
# Using the MIP Map

- In level, find texel and
  - Return the texture value: point sampling (but still better)!
  - Bilinear interpolation
  - Trilinear interpolation



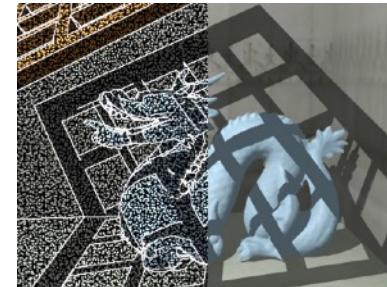
# Using the MIP Map

- In level, find texel and
  - Return the texture value: point sampling (but still better)!
  - Bilinear interpolation
  - Trilinear interpolation



# Memory Usage

- What happens to size of texture?

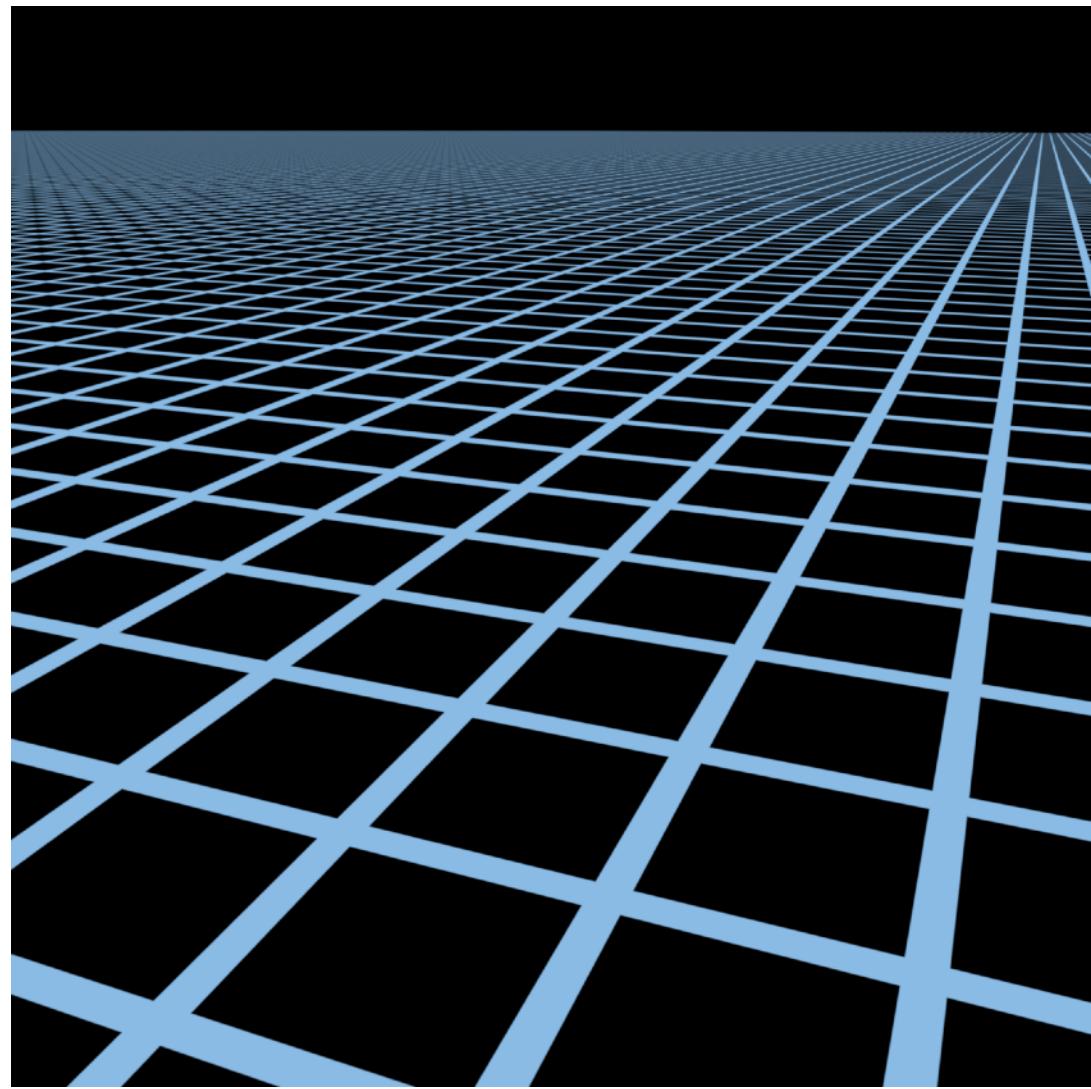


[Kavita Bala]

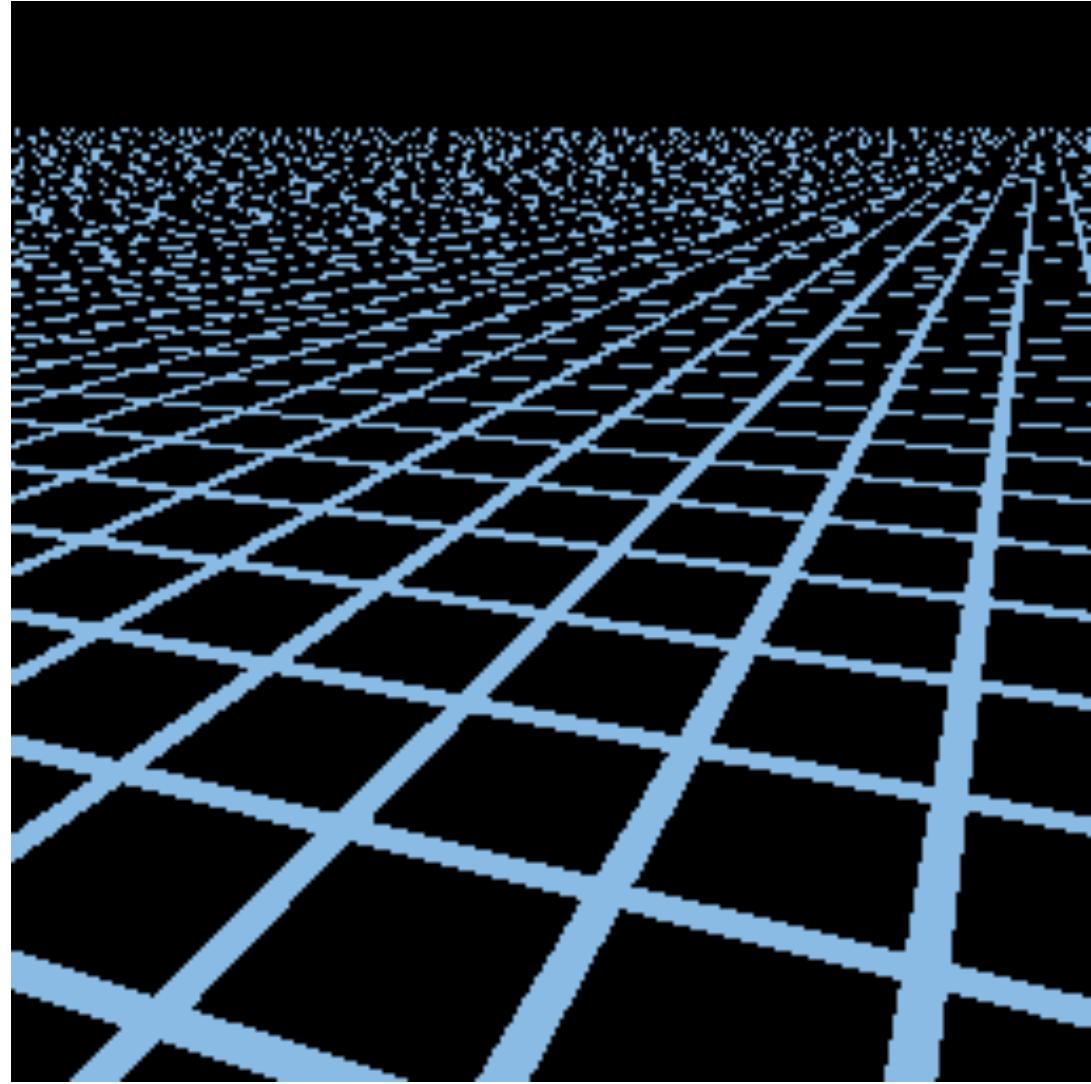
# MIPMAP

- Multi-resolution image pyramid
  - Pre-sampled computation of MIPMAP
  - 1/3 more memory
- Bilinear or Trilinear interpolation

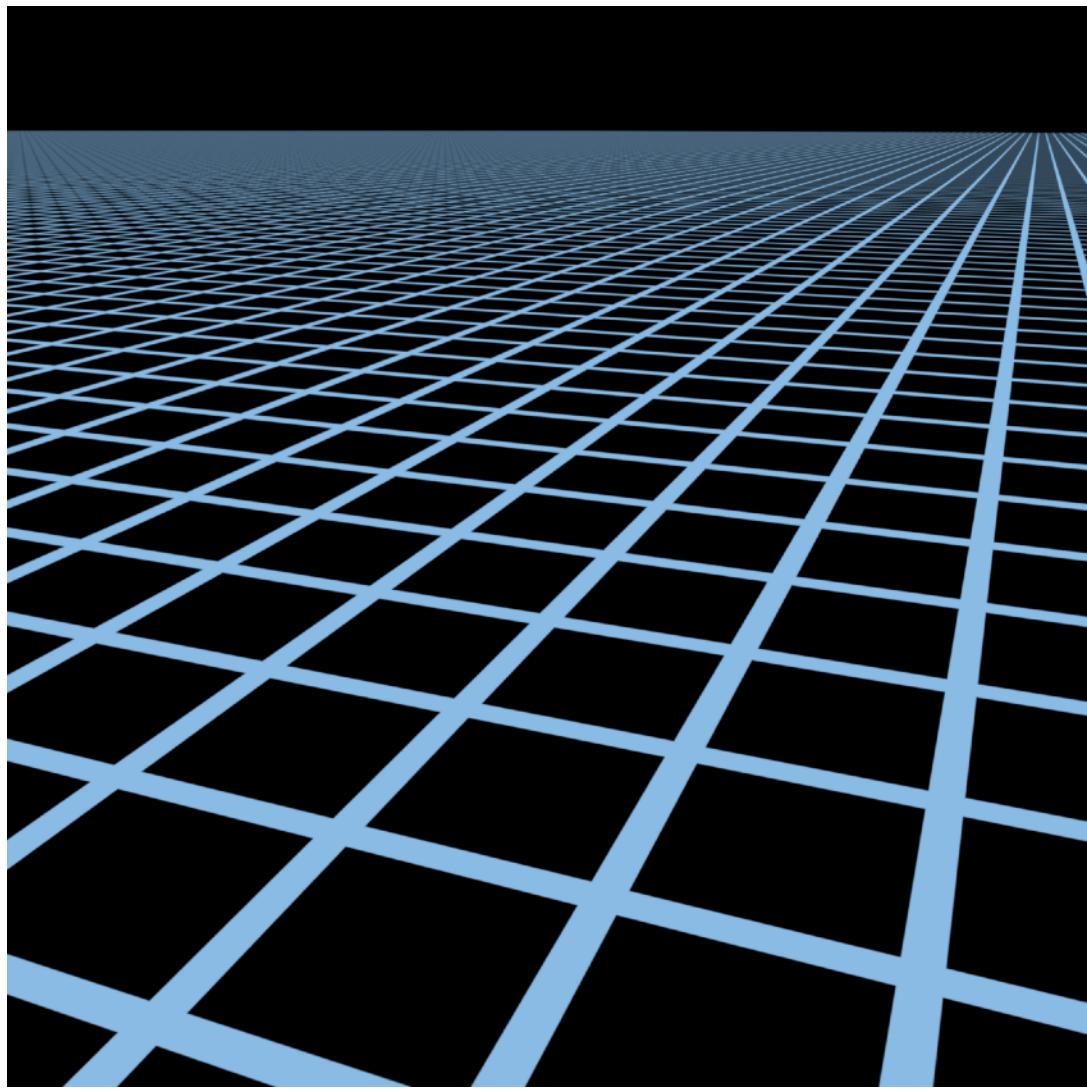
# Point sampling



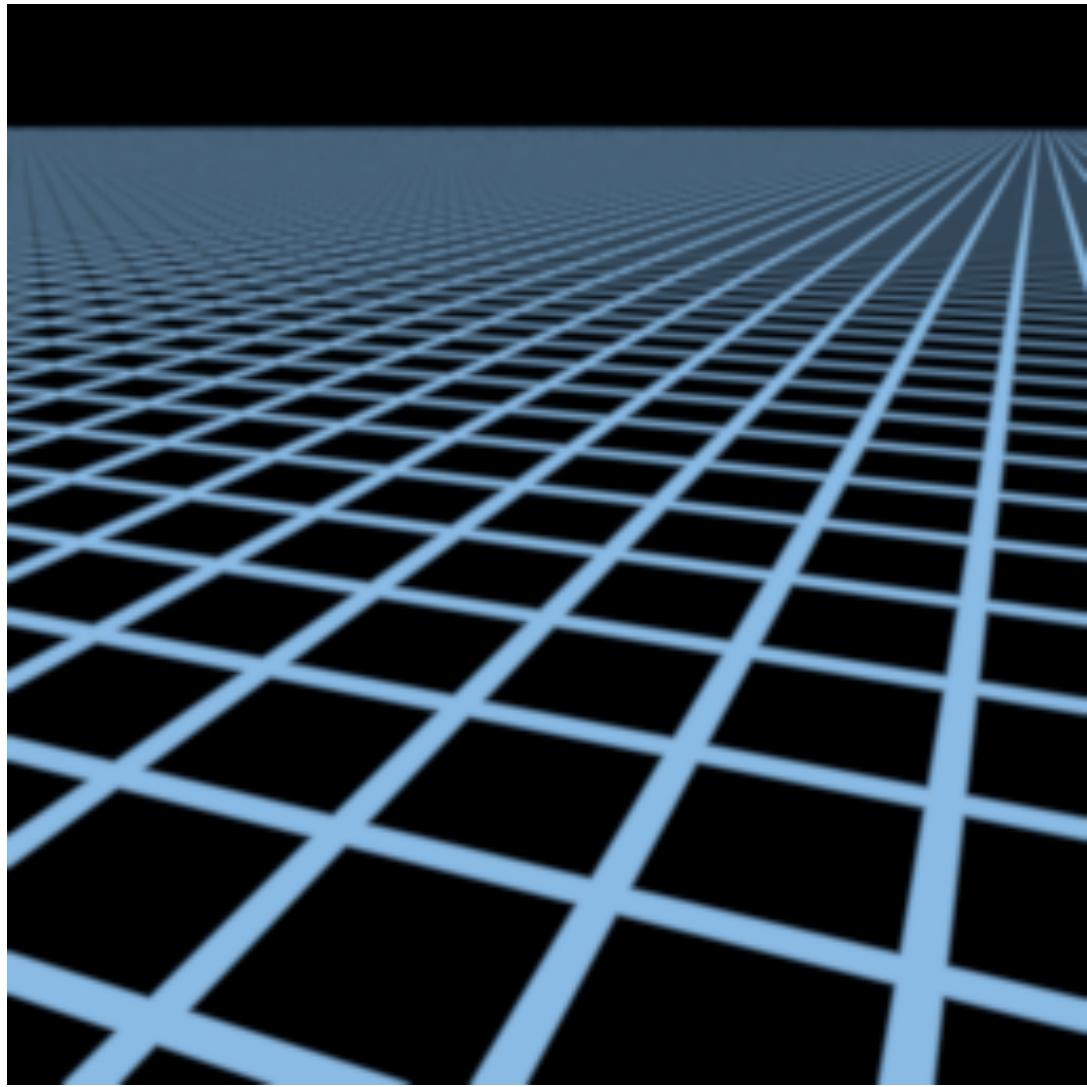
# Point sampling



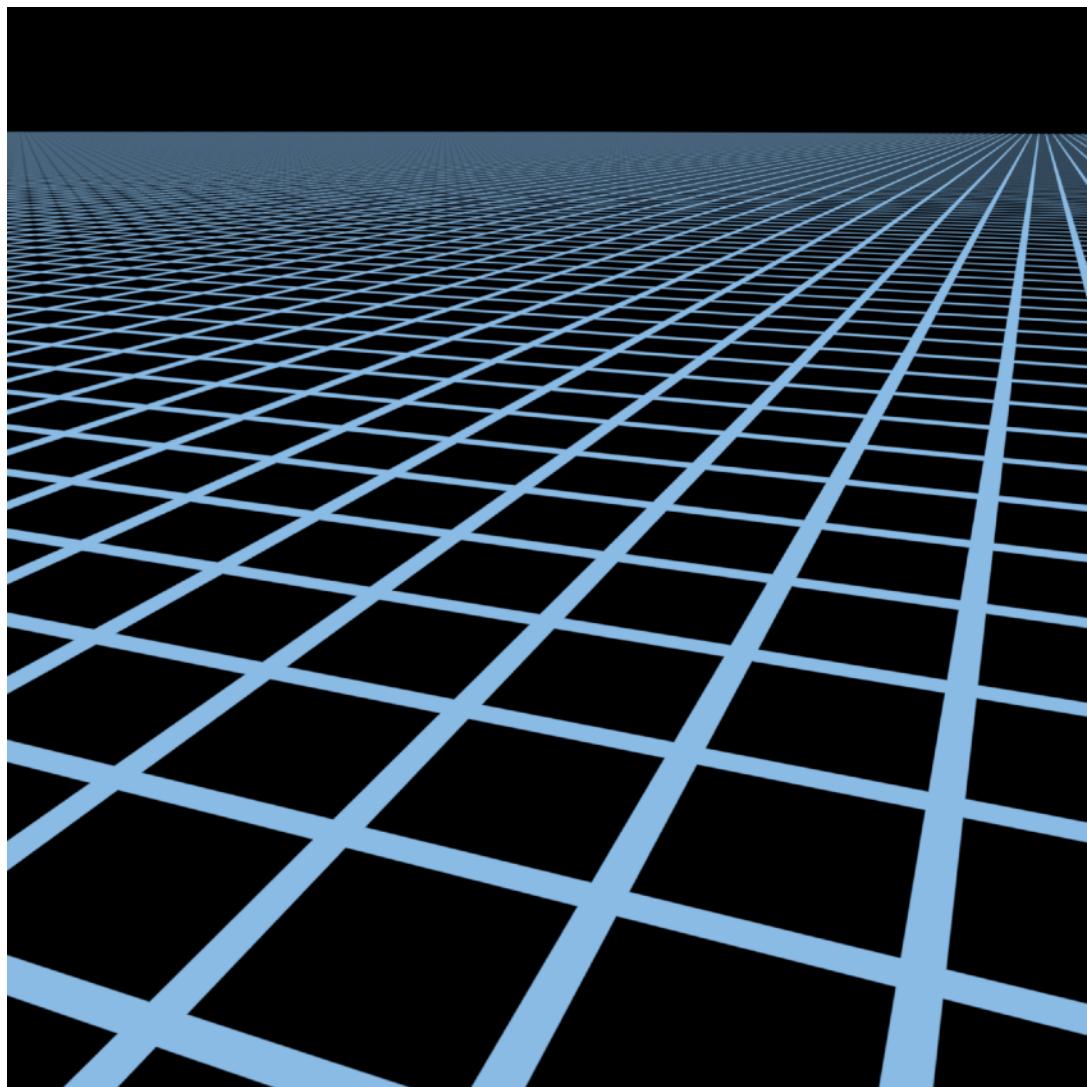
# Reference: gaussian sampling by



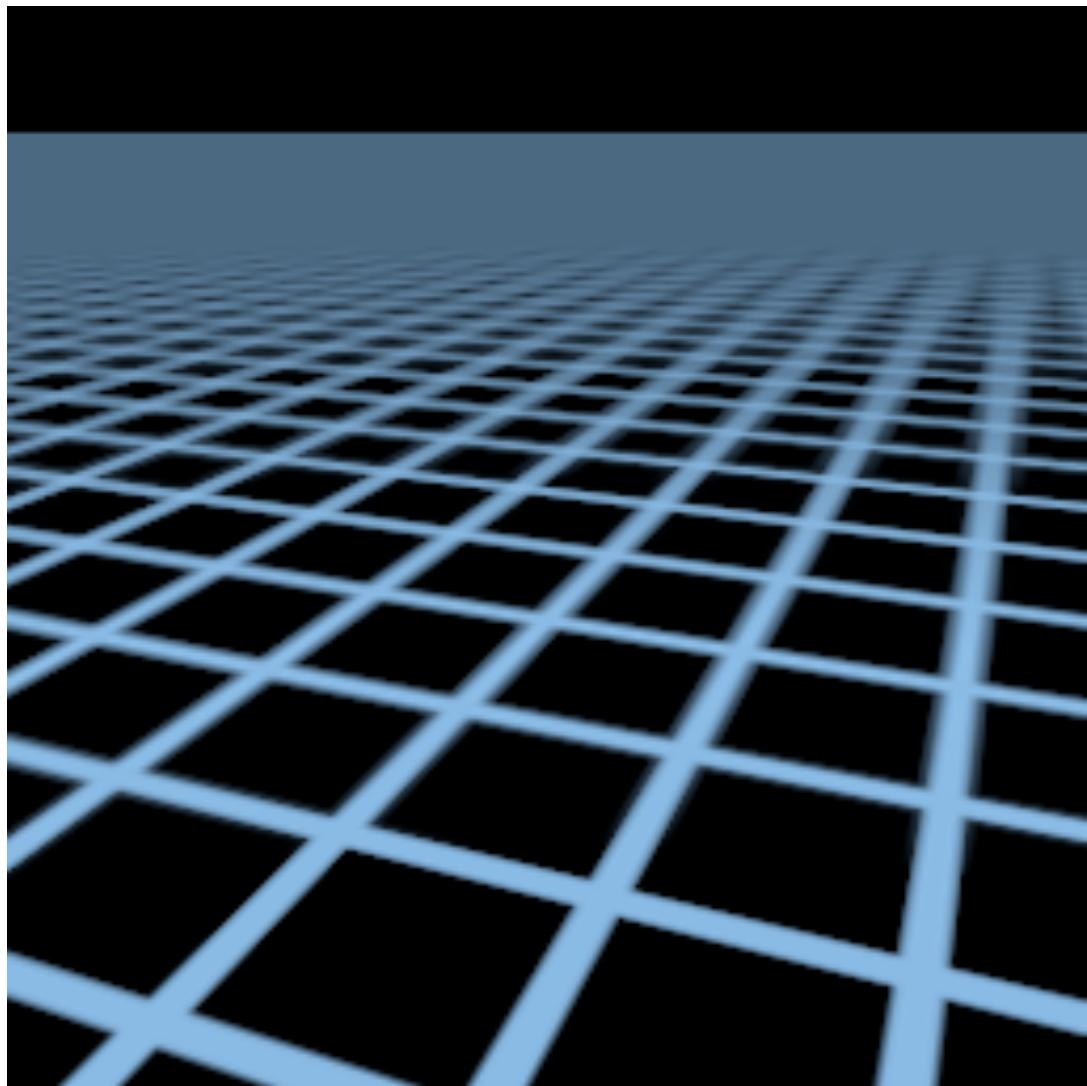
# Reference: gaussian sampling by



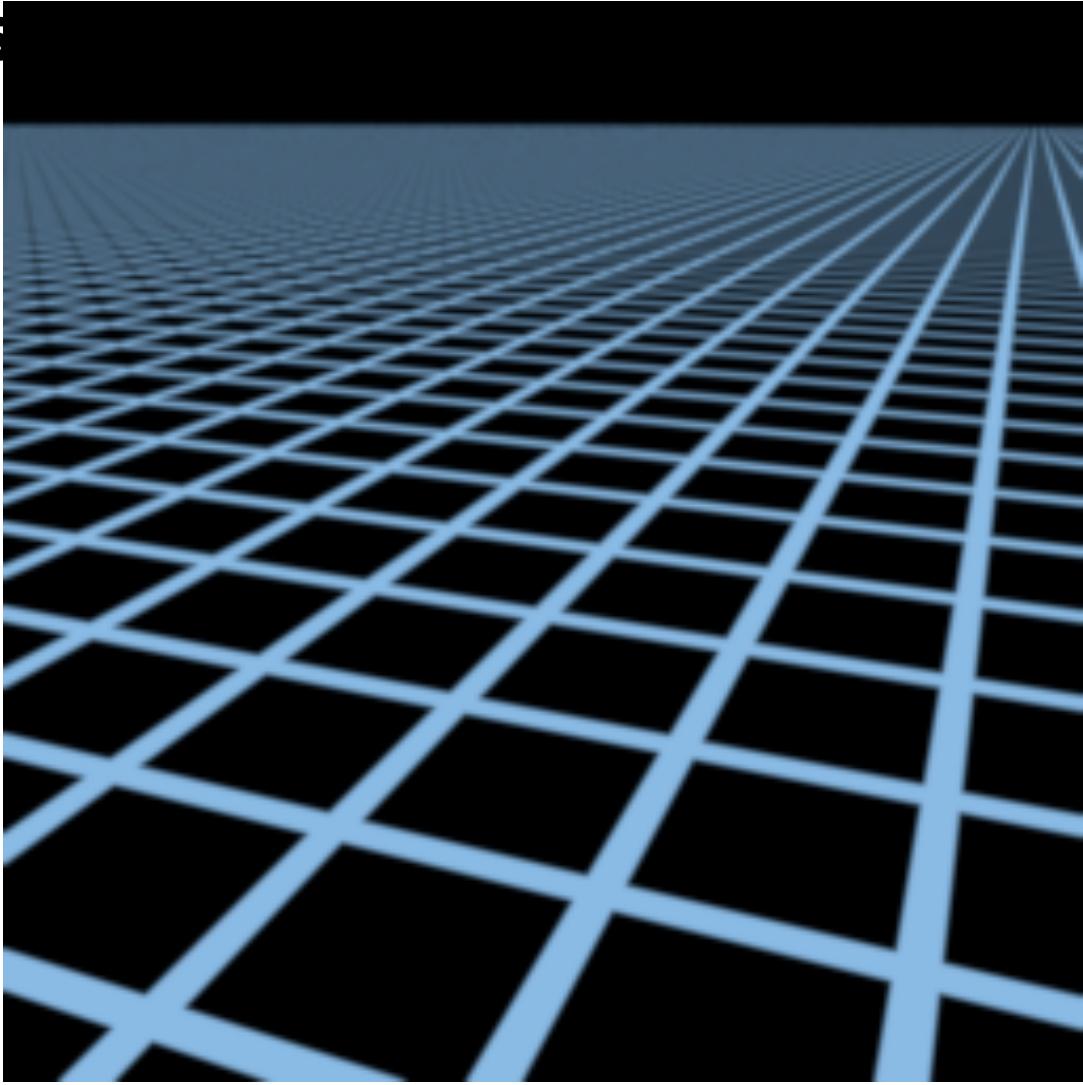
# Texture minification with a mipmap



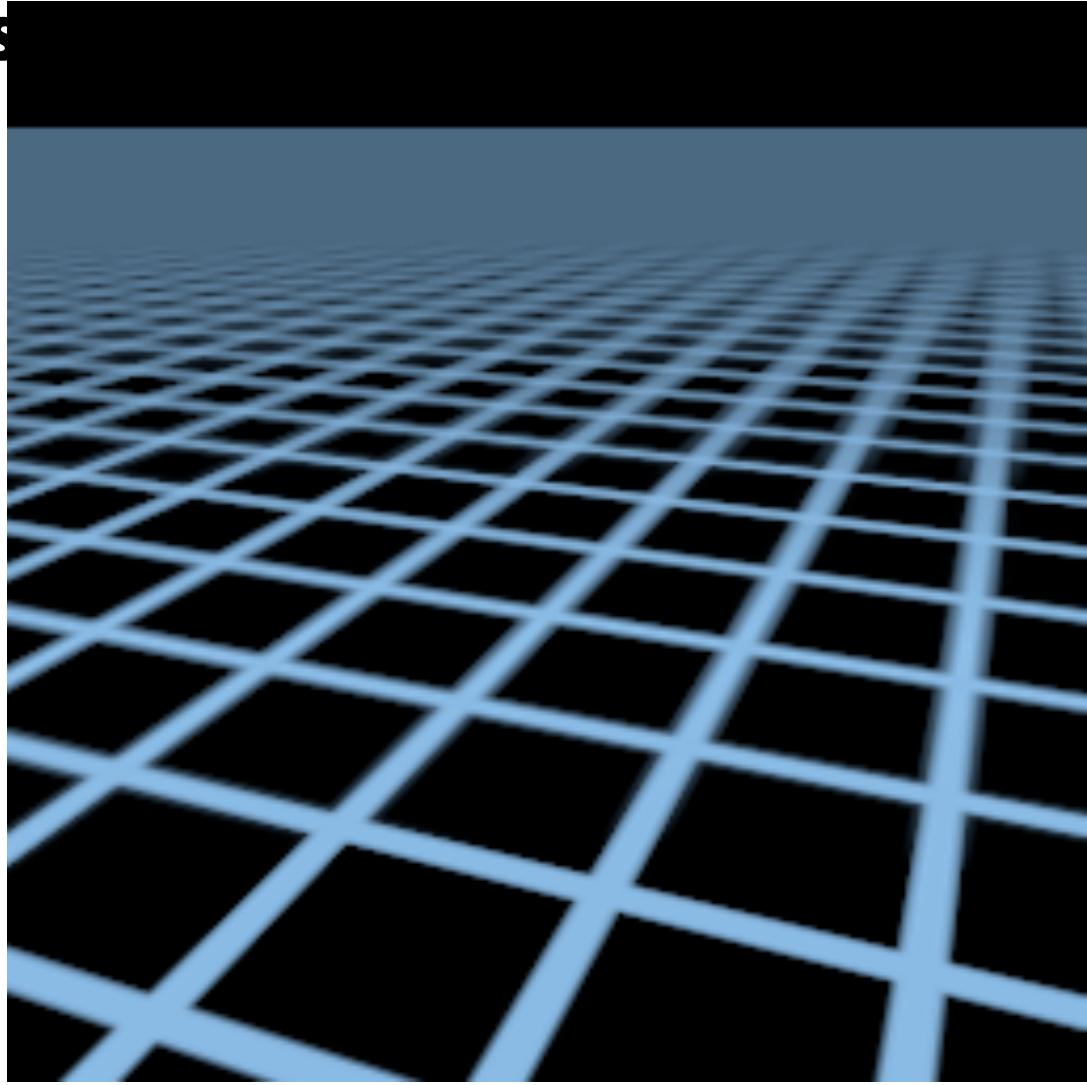
# Texture minification with a mipmap



# Texture minification: supersampling vs.

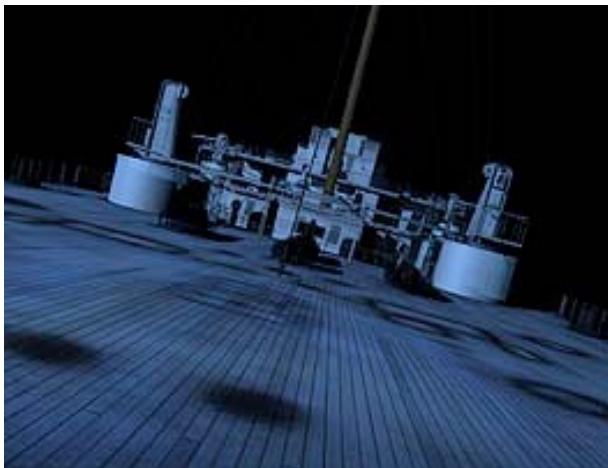
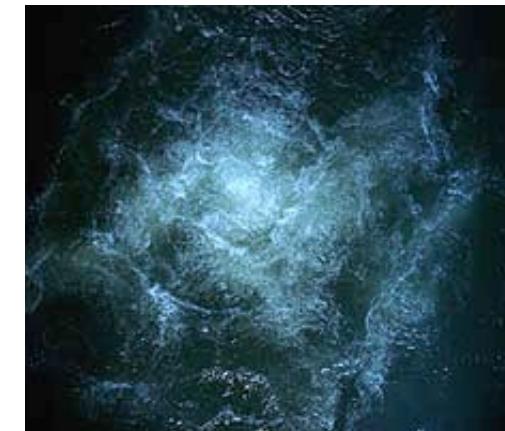


# Texture minification: supersampling vs.



# **Compositing**

# Compositing



[Titanic ; DigitalDomain; vfxhq.com]

# Combining images

- Often useful combine elements of several images
- Trivial example: video crossfade
  - smooth transition from one scene to another

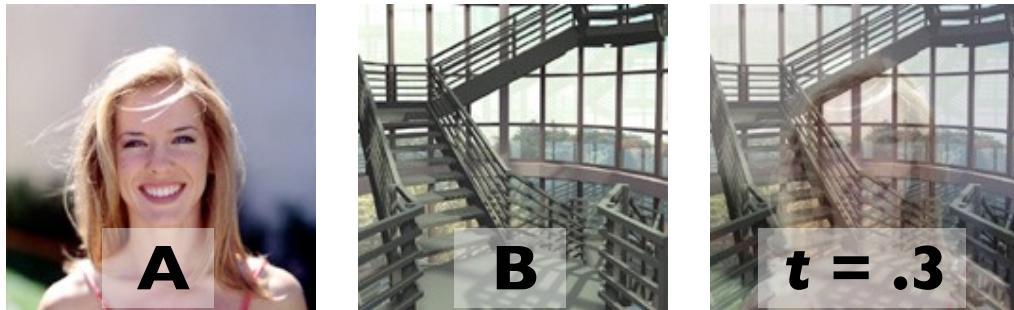


- note: weights sum to 1.0
  - no unexpected brightening or darkening
  - no out-of-range results
- this is *linear interpolation*

$$\begin{aligned}r_C &= tr_A + (1 - t)r_B \\g_C &= tg_A + (1 - t)g_B \\b_C &= tb_A + (1 - t)b_B\end{aligned}$$

# Combining images

- Often useful combine elements of several images
- Trivial example: video crossfade
  - smooth transition from one scene to another



- note: weights sum to 1.0
  - no unexpected brightening or darkening
  - no out-of-range results
- this is *linear interpolation*

$$r_C = tr_A + (1 - t)r_B$$

$$g_C = tg_A + (1 - t)g_B$$

$$b_C = tb_A + (1 - t)b_B$$

# Combining images

- Often useful combine elements of several images
- Trivial example: video crossfade
  - smooth transition from one scene to another



- note: weights sum to 1.0
  - no unexpected brightening or darkening
  - no out-of-range results
- this is *linear interpolation*

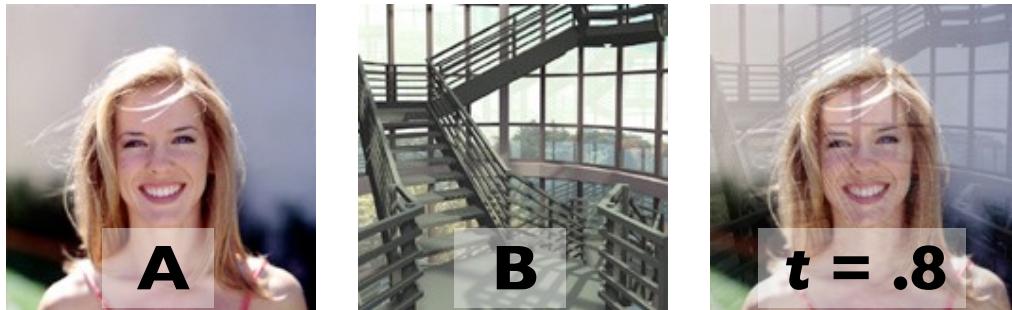
$$r_C = tr_A + (1 - t)r_B$$

$$g_C = tg_A + (1 - t)g_B$$

$$b_C = tb_A + (1 - t)b_B$$

# Combining images

- Often useful combine elements of several images
- Trivial example: video crossfade
  - smooth transition from one scene to another



- note: weights sum to 1.0
  - no unexpected brightening or darkening
  - no out-of-range results
- this is *linear interpolation*

$$\begin{aligned}r_C &= tr_A + (1 - t)r_B \\g_C &= tg_A + (1 - t)g_B \\b_C &= tb_A + (1 - t)b_B\end{aligned}$$

# Combining images

- Often useful combine elements of several images
- Trivial example: video crossfade
  - smooth transition from one scene to another



$$r_C = tr_A + (1 - t)r_B$$

$$g_C = tg_A + (1 - t)g_B$$

$$b_C = tb_A + (1 - t)b_B$$

- note: weights sum to 1.0
  - no unexpected brightening or darkening
  - no out-of-range results
- this is *linear interpolation*

# Foreground and background

- In many cases just adding is not enough
- Example: compositing in film production
  - shoot foreground and background separately
  - also include CG elements
  - this kind of thing has been done in analog for decades
  - how should we do it digitally?

# Foreground and background

- How we compute new image varies with position

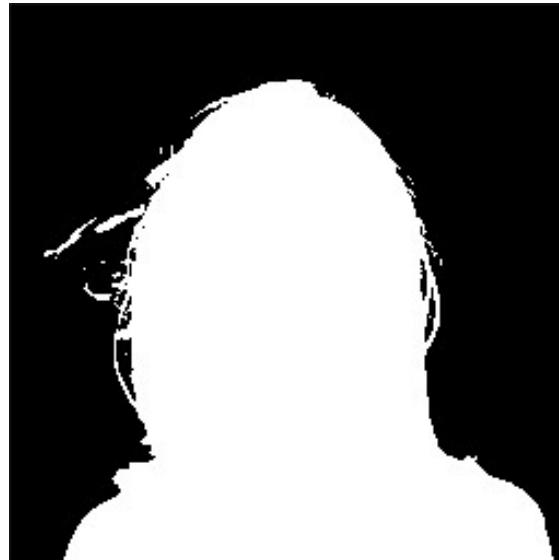


[Chuang et al./Corel]

- Therefore, need to store some kind of tag to say what parts of the image are of interest

# Binary image mask

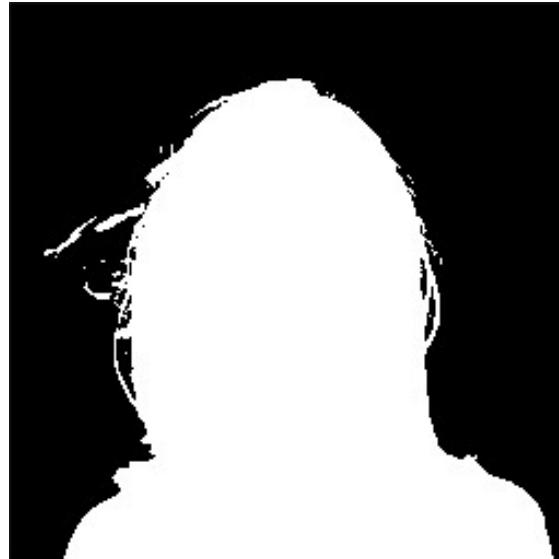
- First idea: store one bit per pixel
  - answers question “is this pixel part of the foreground?”



- causes jaggies similar to point-sampled rasterization
- same problem, same solution: intermediate values

# Binary image mask

- First idea: store one bit per pixel
  - answers question “is this pixel part of the foreground?”



- causes jaggies similar to point-sampled rasterization
- same problem, same solution: intermediate values

# Binary image mask

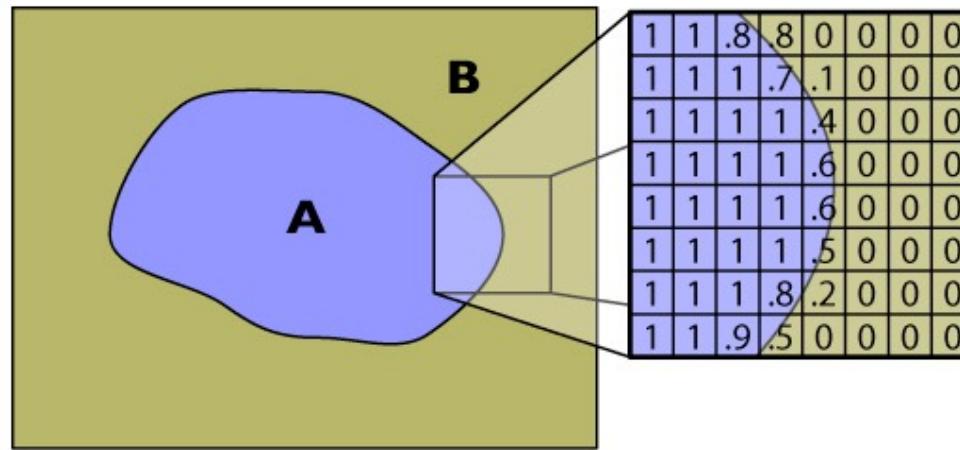
- First idea: store one bit per pixel
  - answers question “is this pixel part of the foreground?”



- causes jaggies similar to point-sampled rasterization
- same problem, same solution: intermediate values

# Partial pixel coverage

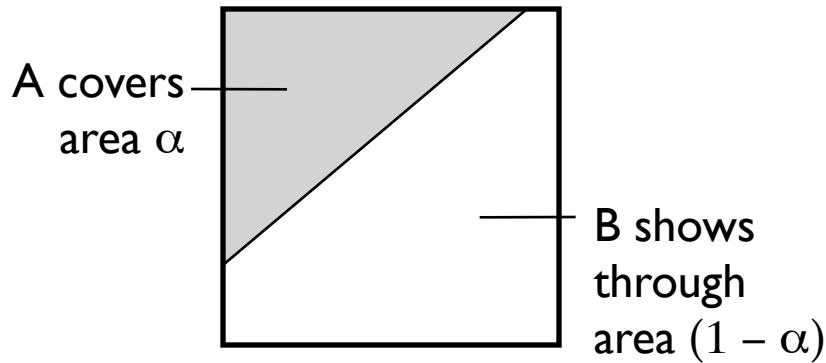
- The problem: pixels near boundary are not strictly foreground or background



- how to represent this simply?
- interpolate boundary pixels between the fg. and bg. colors

# Alpha compositing

- Formalized in 1984 by Porter & Duff
- Store fraction of pixel covered, called  $\alpha$



$$E = A \text{ over } B$$

$$r_E = \alpha_A r_A + (1 - \alpha_A) r_B$$

$$g_E = \alpha_A g_A + (1 - \alpha_A) g_B$$

$$b_E = \alpha_A b_A + (1 - \alpha_A) b_B$$

- this exactly like a spatially varying crossfade
- Convenient implementation
  - 8 more bits makes 32
  - 2 multiplies + 1 add per pixel for compositing

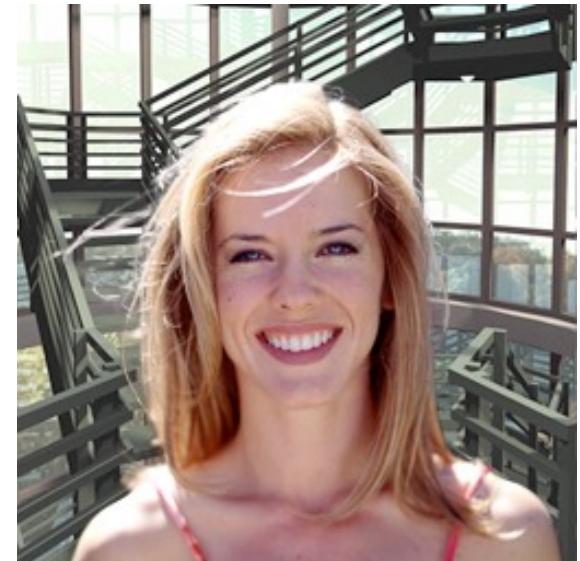
# Alpha compositing—example

[Chuang et al. / Corel]



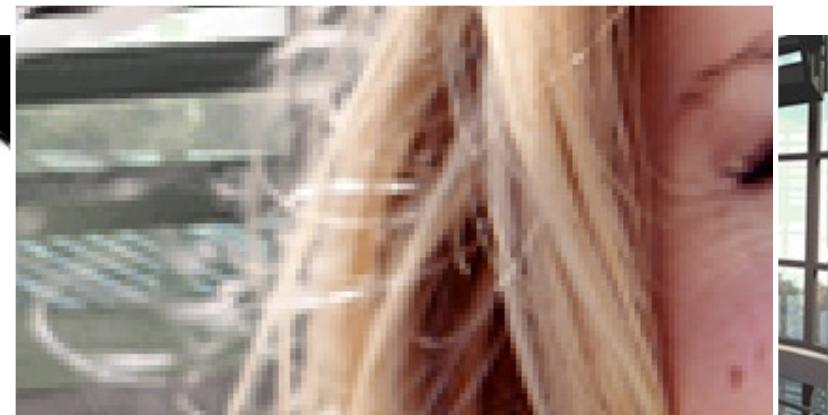
# Alpha compositing—example

[Chuang et al. / Corel]



# Alpha compositing—example

[Chuang et al. / Corel]

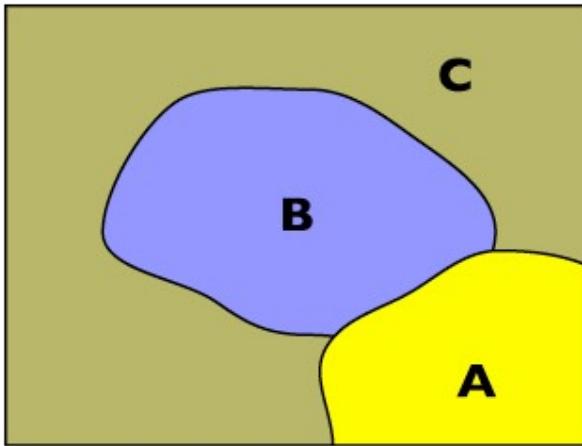


# Compositing composites

- so far have only considered single fg. over single bg.
- in real applications we have  $n$  layers
  - *Titanic* example
  - compositing foregrounds to create new foregrounds
    - what to do with  $\alpha$ ?
- desirable property: associativity
$$A \text{ over } (B \text{ over } C) = (A \text{ over } B) \text{ over } C$$
  - to make this work we need to be careful about how  $\alpha$  is computed

# Compositing composites

- Some pixels are partly covered in more than one layer

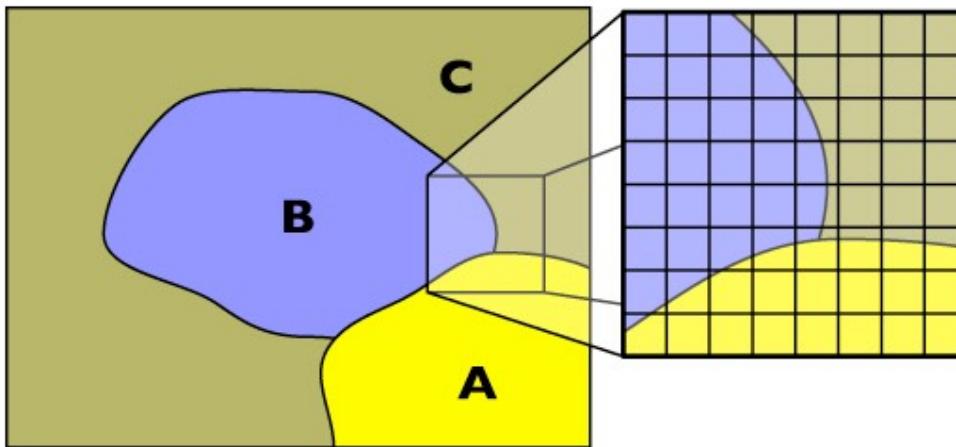


- in  $D = A \text{ over } (B \text{ over } C)$  what will be the result?

$$\begin{aligned}c_D &= \alpha_A c_A + (1 - \alpha_A)[\alpha_B c_B + (1 - \alpha_B)c_C] \\&= \alpha_A c_A + (1 - \alpha_A)\alpha_B c_B + (1 - \alpha_A)(1 - \alpha_B)c_C\end{aligned}$$

# Compositing composites

- Some pixels are partly covered in more than one layer

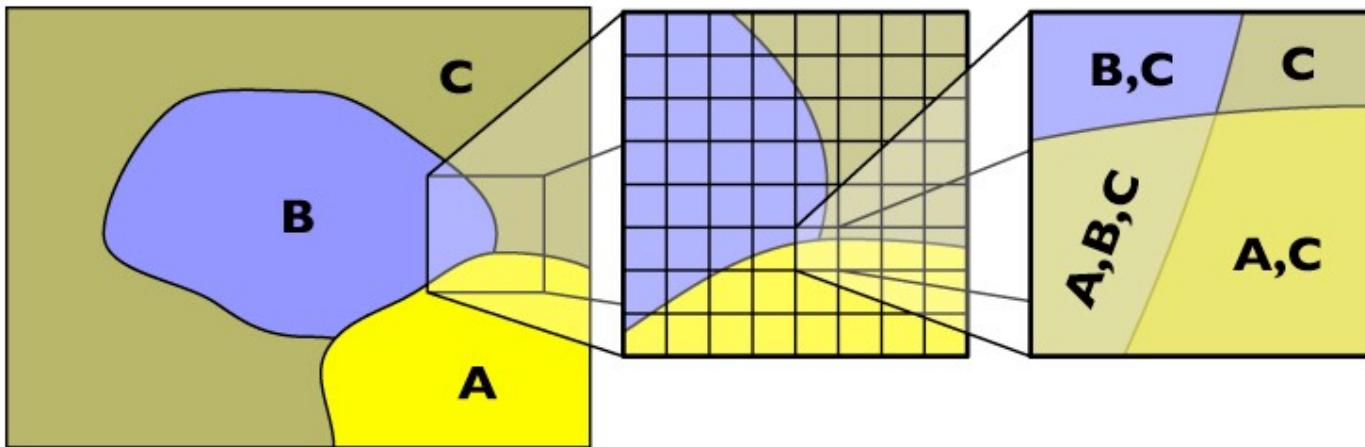


- in  $D = A \text{ over } (B \text{ over } C)$  what will be the result?

$$\begin{aligned}c_D &= \alpha_A c_A + (1 - \alpha_A)[\alpha_B c_B + (1 - \alpha_B)c_C] \\&= \alpha_A c_A + (1 - \alpha_A)\alpha_B c_B + (1 - \alpha_A)(1 - \alpha_B)c_C\end{aligned}$$

# Compositing composites

- Some pixels are partly covered in more than one layer

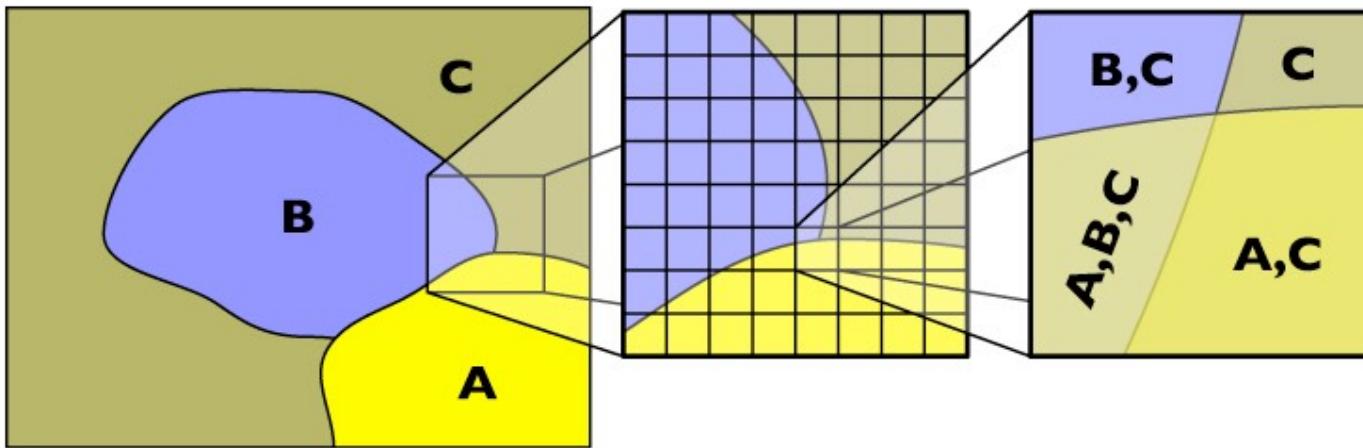


- in  $D = A \text{ over } (B \text{ over } C)$  what will be the result?

$$\begin{aligned}c_D &= \alpha_A c_A + (1 - \alpha_A)[\alpha_B c_B + (1 - \alpha_B)c_C] \\&= \alpha_A c_A + (1 - \alpha_A)\alpha_B c_B + (1 - \alpha_A)(1 - \alpha_B)c_C\end{aligned}$$

# Compositing composites

- Some pixels are partly covered in more than one layer



- in  $D = A \text{ over } (B \text{ over } C)$  what will be the result?

$$\begin{aligned}c_D &= \alpha_A c_A + (1 - \alpha_A)[\alpha_B c_B + (1 - \alpha_B)c_C] \\&= \alpha_A c_A + (1 - \alpha_A)\alpha_B c_B + (1 - \alpha_A)(1 - \alpha_B)c_C\end{aligned}$$

Fraction covered by neither A nor B

# Associativity?

- What does this imply about (A **over** B)?
  - Coverage has to be

$$\begin{aligned}\alpha_{(A \text{ over } B)} &= 1 - (1 - \alpha_A)(1 - \alpha_B) \\ &= \alpha_A + (1 - \alpha_A)\alpha_B\end{aligned}$$

- ...but the color values then don't come out nicely in  $D = (A \text{ over } B) \text{ over } C$ :

$$\begin{aligned}c_D &= \alpha_A c_A + (1 - \alpha_A)\alpha_B c_B + (1 - \alpha_A)(1 - \alpha_B)c_C \\ &= \alpha_{(A \text{ over } B)}(\dots) + (1 - \alpha_{(A \text{ over } B)})c_C\end{aligned}$$

# An optimization

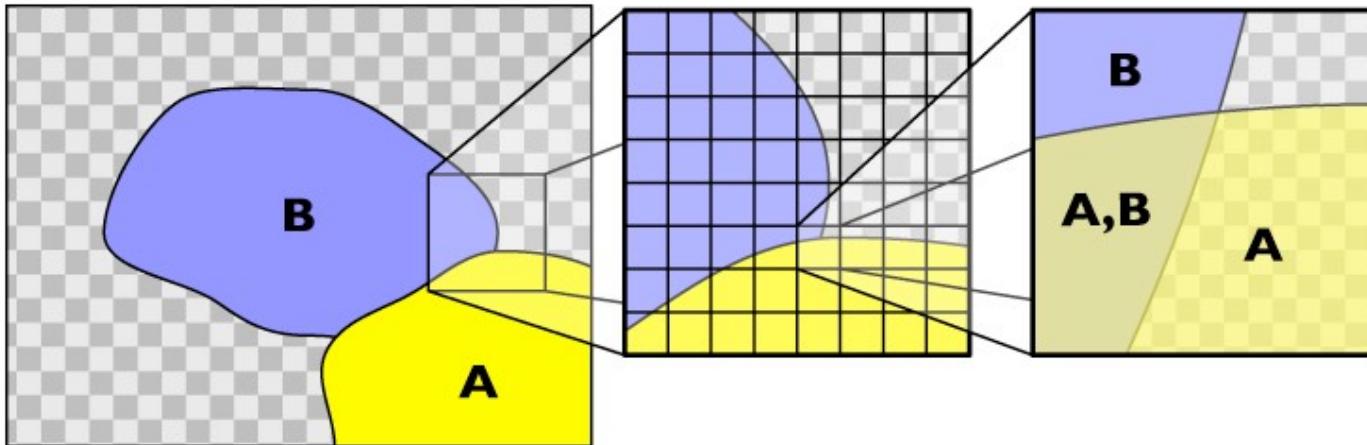
- Compositing equation again

$$c_E = \alpha_A c_A + (1 - \alpha_A) c_B$$

- Note  $c_A$  appears only in the product  $\alpha_A c_A$ 
  - so why not do the multiplication ahead of time?
- Leads to *premultiplied alpha*:
  - store pixel value  $(r', g', b', \alpha)$  where  $c' = \alpha c$
  - $E = A \text{ over } B$  becomes
    - this turns out to be more than an optimization...
    - hint: so far the background has been opaque!

# Compositing composites

- What about just  $E = A \text{ over } B$  (with B transparent)?

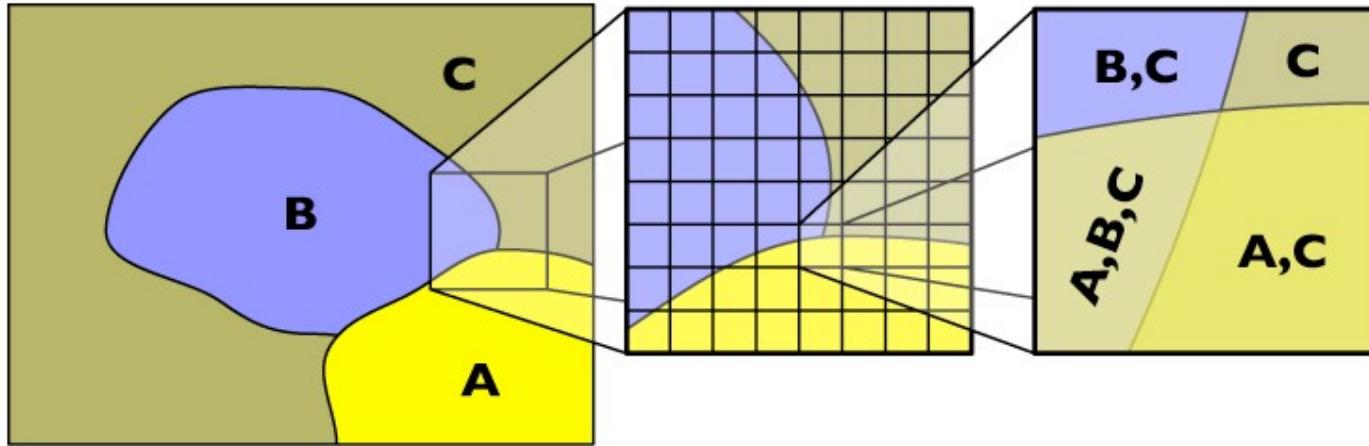


- in premultiplied alpha, the result

$$\alpha_E = \alpha_A + (1 - \alpha_A)\alpha_B$$

looks just like blending colors, and it leads to associativity.

# Associativity!

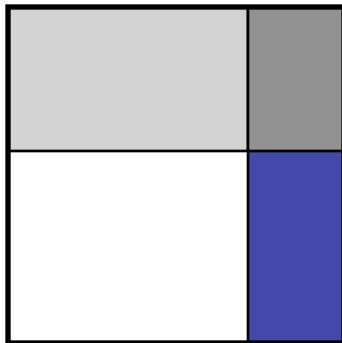


$$\begin{aligned}c_D &= c'_A + (1 - \alpha_A)[c'_B + (1 - \alpha_B)c'_C] \\&= [c'_A + (1 - \alpha_A)c'_B] + (1 - \alpha_A)(1 - \alpha_B)c'_C \\&= c'_{(A \text{ over } B)} + (1 - \alpha_{(A \text{ over } B)})c'_C\end{aligned}$$

- This is another good reason to premultiply

# Independent coverage assumption

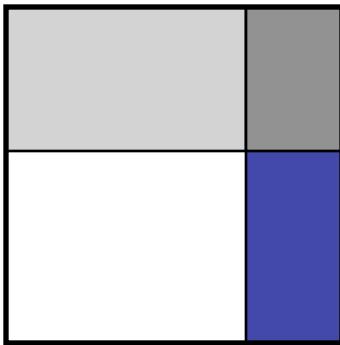
- Why is it reasonable to blend  $\alpha$  like a color?
- Simplifying assumption: covered areas are independent
  - that is, uncorrelated in the statistical sense



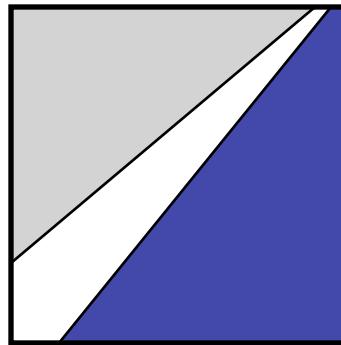
<i>description</i>	<i>area</i>
$\bar{A} \cap \bar{B}$	$(1-\alpha_A)(1-\alpha_B)$
$A \cap \bar{B}$	$\alpha_A(1-\alpha_B)$
$\bar{A} \cap B$	$(1-\alpha_A)\alpha_B$
$A \cap B$	$\alpha_A\alpha_B$

# Independent coverage assumption

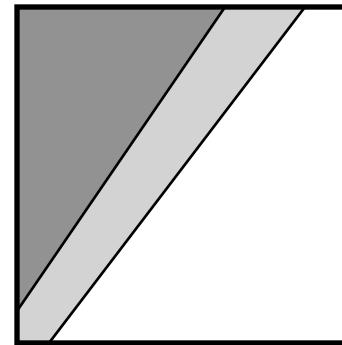
- Holds in most but not all cases



this



not this



or this

- This will cause artifacts
  - but we'll carry on anyway because it is simple and usually works...

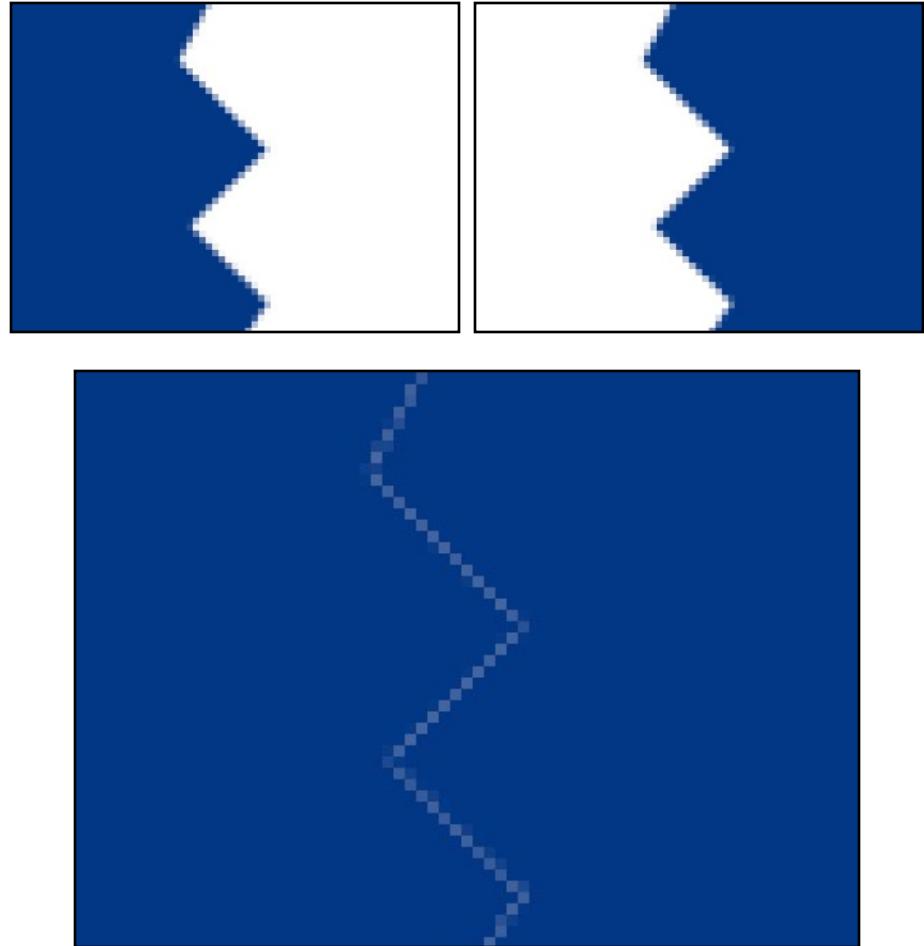
# Alpha compositing—failures

[Cornell PCG]



[Chuang et al. / Core]

positive correlation:  
too much foreground



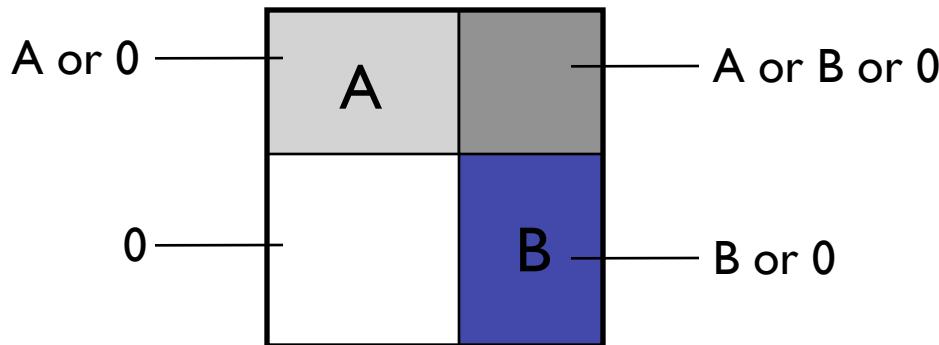
negative correlation:  
too little foreground

# Other compositing operations

- Generalized form of compositing equation:

$$\alpha_E = A \text{ op } B$$

$$c'_E = F_A c'_A + F_B c'_B$$



$1 \times 2 \times 3 \times 2 = 12$  reasonable choices

operation	quadruple	diagram	$F_A$	$F_B$
clear	(0,0,0,0)		0	0
$A$	(0,A,0,A)		1	0
$B$	(0,0,B,B)		0	1
$A \text{ over } B$	(0,A,B,A)		1	$1-\alpha_A$
$B \text{ over } A$	(0,A,B,B)		$1-\alpha_B$	1
$A \text{ in } B$	(0,0,0,A)		$\alpha_B$	0
$B \text{ in } A$	(0,0,0,B)		0	$\alpha_A$
$A \text{ out } B$	(0,A,0,0)		$1-\alpha_B$	0
$B \text{ out } A$	(0,0,B,0)		0	$1-\alpha_A$
$A \text{ atop } B$	(0,0,B,A)		$\alpha_B$	$1-\alpha_A$
$B \text{ atop } A$	(0,A,0,B)		$1-\alpha_B$	$\alpha_A$
$A \text{ xor } B$	(0,A,B,0)		$1-\alpha_B$	$1-\alpha_A$